


ICT1010 Computer Networks

Lab Exercise – TCP



Objective

To understand the working of the reliable TCP (Transmission Control Protocol). TCP is the main transport layer protocol used in the Internet.

Step 1: Open the Trace

Download, Unzip and Open the trace file from your LMS:

<https://xsite.singaporetech.edu.sg/d2l/le/content/73233/viewContent/458862/View>

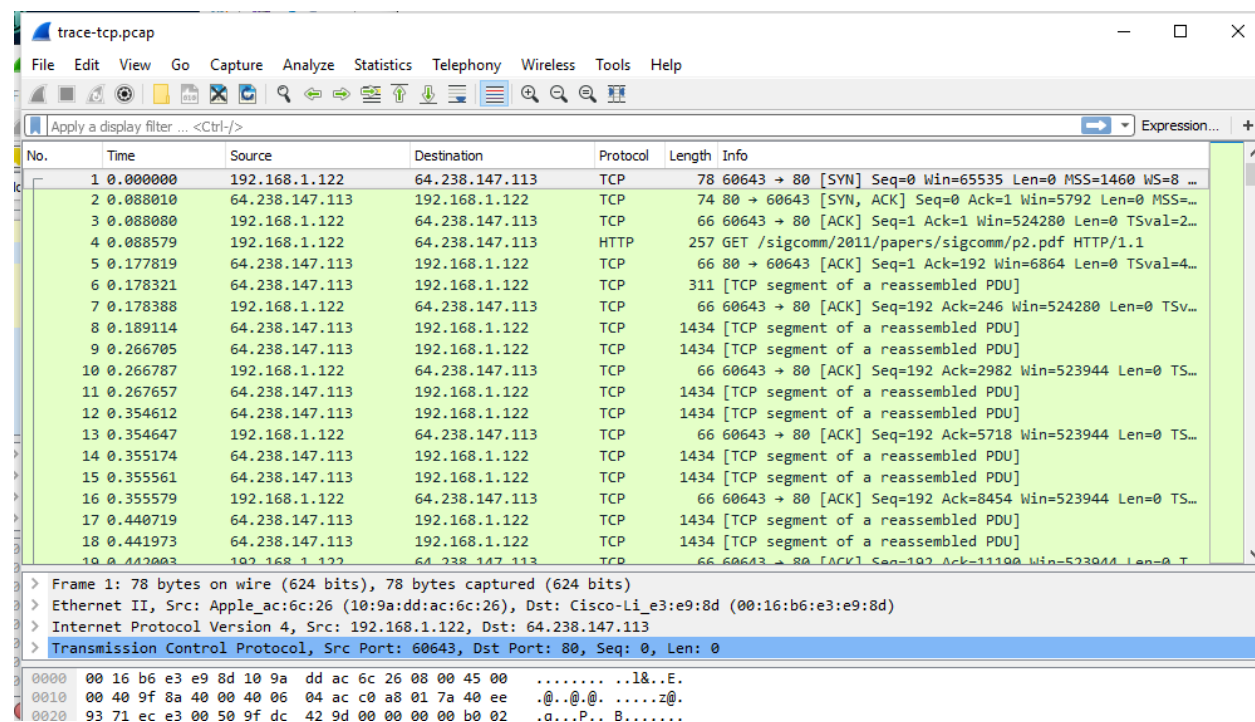


Figure 1: Open the trace-tcp.pcap file

Step 2: Inspect the Trace

1. Select a **long packet** anywhere in the middle of your trace whose protocol is listed as TCP. See below in Figure 2., where a TCP packet of length 66 is highlighted. All packets except the initial HTTP GET and last packet of the HTTP response should be listed as TCP. Picking a long packet ensures that we are looking at a download packet from the server to a computer.
2. Looking at the protocol layers, you should see an IP block before the TCP block. This is because the TCP segment is carried in an IP. xpend the TCP protocol section by using the ">" expander or icon.

You will see roughly the following fields:

- First is the source port, then the destination port. This is the addressing that TCP adds beyond the IP address.

What is the source port number? 80

Whats does the above source port number signify? http

- Then there is the sequence number field.

What is the need for a sequence number? It is to ensure that missing or wrongly ordered packets can be detected and fixed

- Next is the acknowledgement field. It tells the last received position in the reverse byte stream.
- The header length giving the length of the TCP header.
- The flags field has multiple flag bits to indicate the type of TCP segment. You can expand it and look at the possible flags.
- Next is a checksum, to detect transmission errors.
- There may be an Options field with various options. You can expand this field and explore if you would like.
- Finally, there may be a TCP payload, carrying the bytes that are being transported.

As well as the above fields, there may be other informational lines that Wireshark provides to help you interpret the packet. We have covered only the fields that are carried across the network.

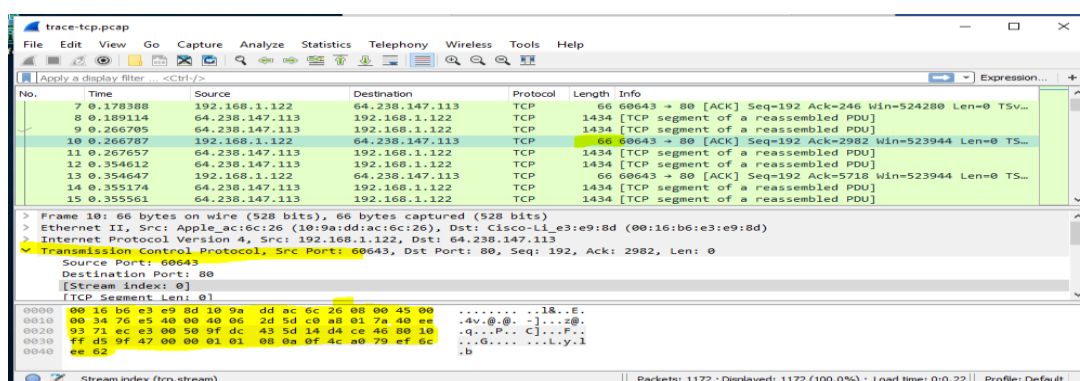


Figure 2: Examining the TCP segment feilds

Step 3: Understanding TCP Segment Structure




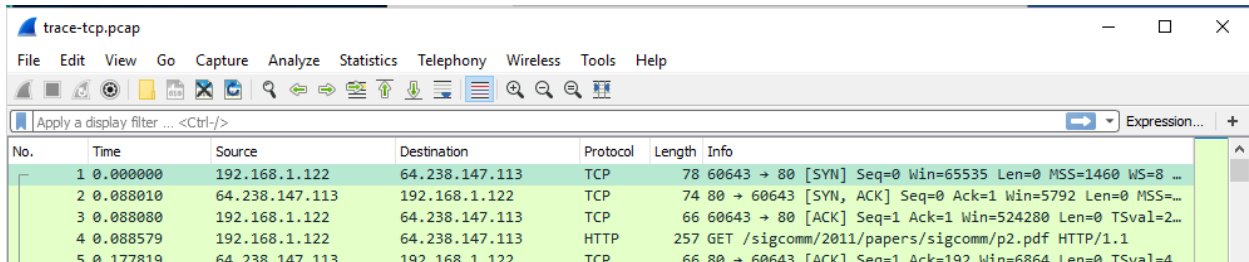
Figure 3: Structure of a TCP segment

- The Header length and Flags fields are combined into a 2-byte quantity. It is not easy to determine their bit lengths with Wireshark.
- The Urgent Pointer field is shown as dotted. This field is typically not used, and so does not show up in Wireshark and we do not expect you to have it in your drawing. You can notice its existence in Wireshark, however, by observing the zero bytes in the segment that are skipped over as you select the different fields.
- The Options field is shown dotted, as it may or may not be present for the segments in your trace. Most often it will be present, and when it is then its length will be a multiple of four bytes.
- The Payload is optional. It is present for the segment you viewed, but not present on an Ack-only segment, for example.

Step 4: TCP Connection Setup/Teardown

Three-Way Handshake

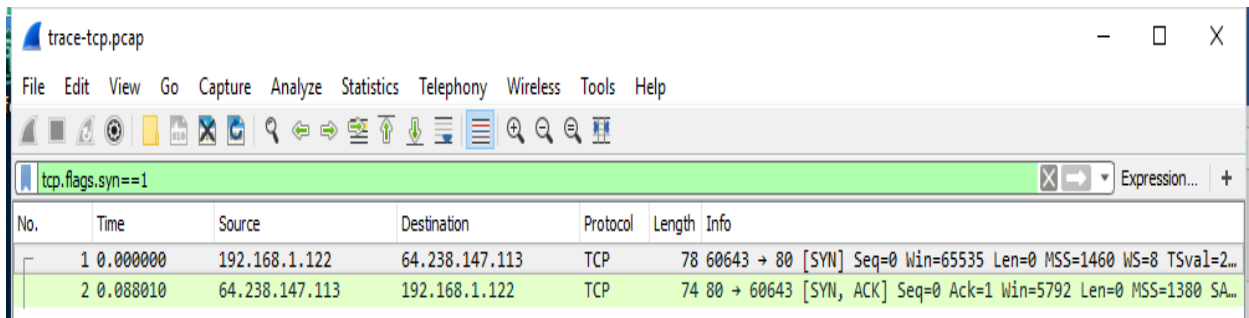
To see the “three way handshake” in action, look for a TCP segment with the SYN flag on. These are up at the beginning of your trace, and the packets that follow it (see below).



No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	192.168.1.122	64.238.147.113	TCP	78	60643 → 80 [SYN] Seq=0 Win=65535 Len=0 MSS=1460 WS=8 ...
2	0.088010	64.238.147.113	192.168.1.122	TCP	74	80 → 60643 [SYN, ACK] Seq=0 Ack=1 Win=5792 Len=0 MSS=...
3	0.088080	192.168.1.122	64.238.147.113	TCP	66	60643 → 80 [ACK] Seq=1 Ack=1 Win=524280 Len=0 TSval=2...
4	0.088579	192.168.1.122	64.238.147.113	HTTP	257	GET /sigcomm/2011/papers/sigcomm/p2.pdf HTTP/1.1
5	0.177819	64.238.147.113	192.168.1.122	TCP	66	80 → 60643 [ACK] Seq=1 Ack=192 Win=6864 Len=0 TSval=4...

Figure 4: Selecting a TCP segment with SYN flag

The SYN flag is noted in the Info column. You can also search for packets with the SYN flag on using the filter expression “`tcp.flags.syn==1`”. (See below)



No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	192.168.1.122	64.238.147.113	TCP	78	60643 → 80 [SYN] Seq=0 Win=65535 Len=0 MSS=1460 WS=8 TSval=2...
2	0.088010	64.238.147.113	192.168.1.122	TCP	74	80 → 60643 [SYN, ACK] Seq=0 Ack=1 Win=5792 Len=0 MSS=1380 SA...

Figure 5: Selecting a TCP segment with SYN flag on

A “SYN packet” is the start of the three-way handshake. In this case it is being sent from your computer to the remote server. The remote server should reply with a TCP segment with the SYN and ACK flags set, or a “SYN ACK packet”. On receiving this segment, your computer will ACK it, consider the connection set up, and begin sending data, which in this case will be the HTTP request.

Step 5a: TCP Connection Setup/Teardown

Next, we wish to clear the display filter `tcp.flags.syn==1` so that we can once again see all the packets in our original trace. Do this by clearing the display filter as shown below.




Figure 6: Clearing the display filter TCP segment with SYN flag on

If you do this correctly, you should see the full trace. We are most interested in the first three packets to understand the three-way handshake.

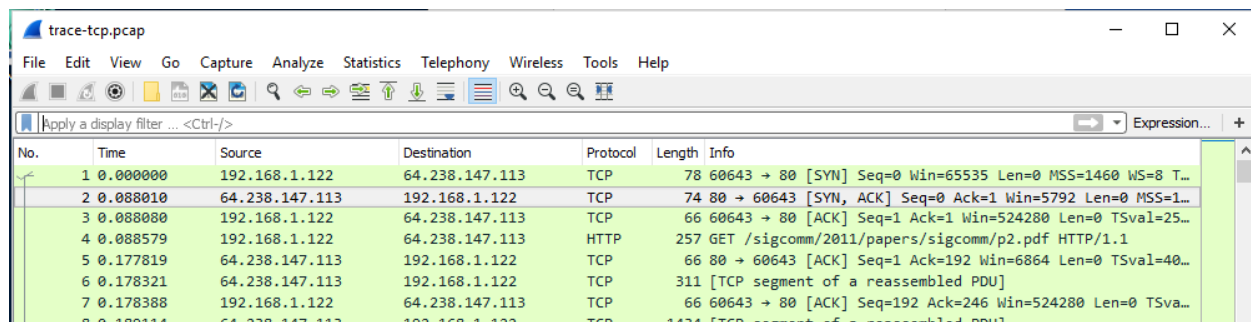


Figure 7: Viewing the complete trace

Below is a time sequence diagram of the three-way handshake in your trace, up to and including the first data packet (the HTTP GET request) sent by 'your computer' when the connection is established. As usual, time runs down the page, and lines across the page indicate segments.

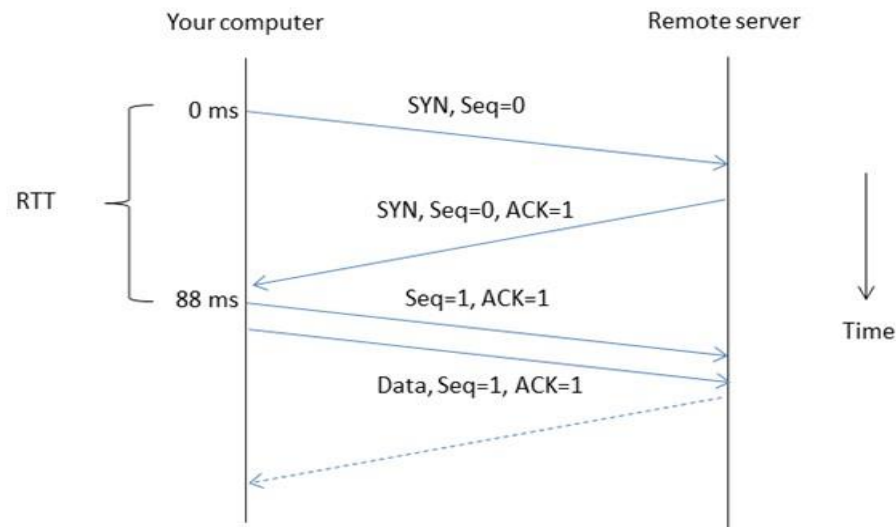


Figure 8: Time sequence diagram for the TCP three-way handshake

There are several features to note:

- The initial SYN has no ACK number, only a sequence number. All subsequent packets have ACK numbers.
- The initial sequence numbers are shown as zero in each direction. This is because Wireshark is configured to show relative sequence numbers. The actual sequence number is some large 32-bit number, and it is different for each end.
- The ACK number in the three-way handshake is the previous segment length plus 1.
- Our computer sends the third part of the handshake (the ACK) and then sends data right away in a different packet. It would be possible to combine these packets, but they are typically separate (because one is triggered by the OS and one by the application).
- For the Data segment, the sequence number and ACK stay with the previous values. The sequence number will advance as the sender sends more data. The ACK number will advance as the sender receives more data from the remote server.
- The time taken for the three packets to be received and sent happen very close together compared to the gap between the first and second packet.

What is the observed RTT in our trace? about 0.02 seconds

- If you use a local web server, the RTT will be very small, likely a few milliseconds. If you use a major web server that may be provided by a content distribution network, the RTT will likely be tens of milliseconds. If you use a geographically remote server, the RTT will likely be hundreds of milliseconds.

Step 5b: Connection Options

As well as setting up a connection, the TCP SYN packets negotiate parameters between the two ends using Options. Each end describes its capabilities, if any, to the other end by including the appropriate Options on its SYN. Often both ends must support the behavior for it to be used during data transfer.

Common Options include Maximum Segment Size (MSS) to tell the other side the largest segment that can be received, and Timestamps to include information on segments for estimating the round trip time. There are also Options such as NOP (No-operation) and End of Option list that serve to format the Options but do not advertise capabilities. You do not need to include these formatting options in your answer above. Options can also be carried on regular segments after the connection is set up when they play a role in data transfer. This depends on the Option. For example: the MSS option is not carried on each packet because it does not convey new information; timestamps may be included on each packet to keep a fresh estimate of the RTT; and options such as SACK (Selective Acknowledgments) are used only when data is received out of order.

Our TCP Options are Maximum Segment Size, Window Scale, SACK permitted, and Timestamps. Each of these Options is used in both directions. There are also the NOP & End of Option List formatting options.

Here is an example of a FIN teardown:

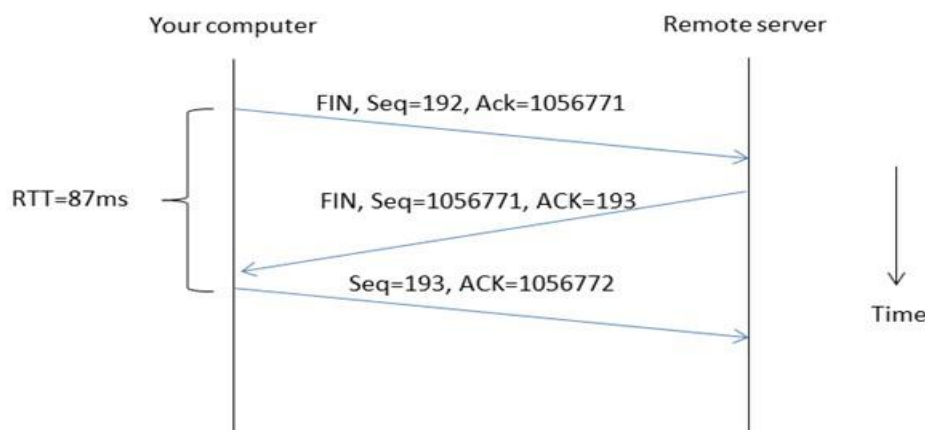


Figure 9: Time sequence diagram for FIN teardown

Points to note:

- The teardown process in Figure.9 is initiated by the computer; it might be initiated by the server as well.
- Like the SYN, the FIN flag occupies one sequence number. Thus, when the sequence number of the FIN is 192, the corresponding Ack number is 193.
- Your sequence numbers will vary. Our numbers are relative (as computed by Wireshark) but clearly depend on the resource that is fetched. You can tell that it is around 1 MB long.
- The RTT in the FIN exchange is like that in the SYN exchange, as it should be. Your RTT will vary depending on the distance between the computer and server as before.

Step 6: FIN/RST Teardown

Finally, the TCP connection is taken down after the download is complete. This is typically done with FIN (Finalize) segments. Each side sends a FIN to the other and acknowledges the FIN they receive; it is similar to the three-way handshake as observed in Figure 9.

Alternatively, the connection may be torn down abruptly when one end sends a RST (Reset). However, this packet does not need to be acknowledged by the other side (see below Figure 10).

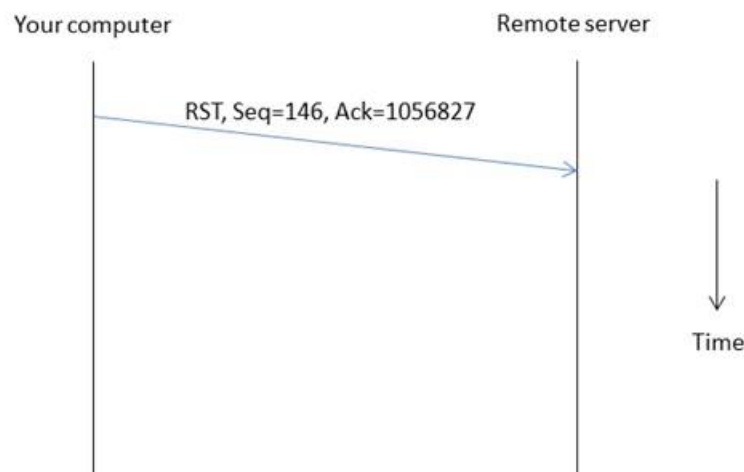
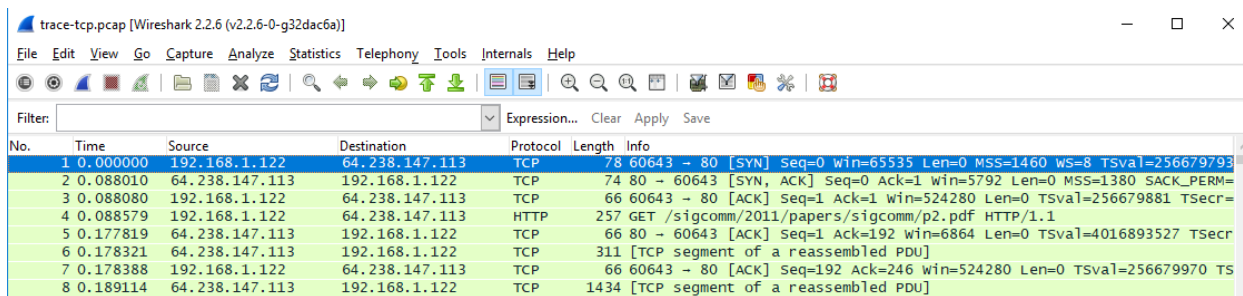


Figure 10: Time sequence diagram for RST teardown

Points to note:

- The teardown is initiated by the computer; it might also be initiated by the server.
- The teardown is abrupt – a single RST in this case, and then it is closed, which the other end must accommodate.
- Since there is no round trip exchange, no RTT can be estimated.

Step 7: TCP Data Transfer



No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	192.168.1.122	64.238.147.113	TCP	78	60643 → 80 [SYN] Seq=0 Win=65535 Len=0 MSS=1460 WS=8 TSval=256679793
2	0.088010	64.238.147.113	192.168.1.122	TCP	74	80 → 60643 [SYN, ACK] Seq=0 Ack=1 win=5792 Len=0 MSS=1380 SACK_PERM=
3	0.088080	192.168.1.122	64.238.147.113	TCP	66	60643 → 80 [ACK] Seq=1 Ack=1 win=524280 Len=0 TSval=256679881 TSecr=
4	0.088579	192.168.1.122	64.238.147.113	HTTP	257	GET /sigcomm/2011/papers/sigcomm/p2.pdf HTTP/1.1
5	0.177819	64.238.147.113	192.168.1.122	TCP	66	80 → 60643 [ACK] Seq=1 Ack=192 win=6864 Len=0 TSval=4016893527 TSecr=
6	0.178321	64.238.147.113	192.168.1.122	TCP	311	[TCP segment of a reassembled PDU]
7	0.178388	192.168.1.122	64.238.147.113	TCP	66	60643 → 80 [ACK] Seq=192 Ack=246 win=524280 Len=0 TSval=256679970 TS
8	0.189114	64.238.147.113	192.168.1.122	TCP	1434	[TCP segment of a reassembled PDU]

The middle portion of the TCP connection is the data transfer, or download, in our trace. This is the main event. To get an overall sense of it, we will first look at the download rate over time.

Under the *Statistics* menu select an “IO Graph” (as shown below).




Figure 11: Opening an IO graph

You should end up with a graph like below. By default, this graph shows the rate of packets over time.

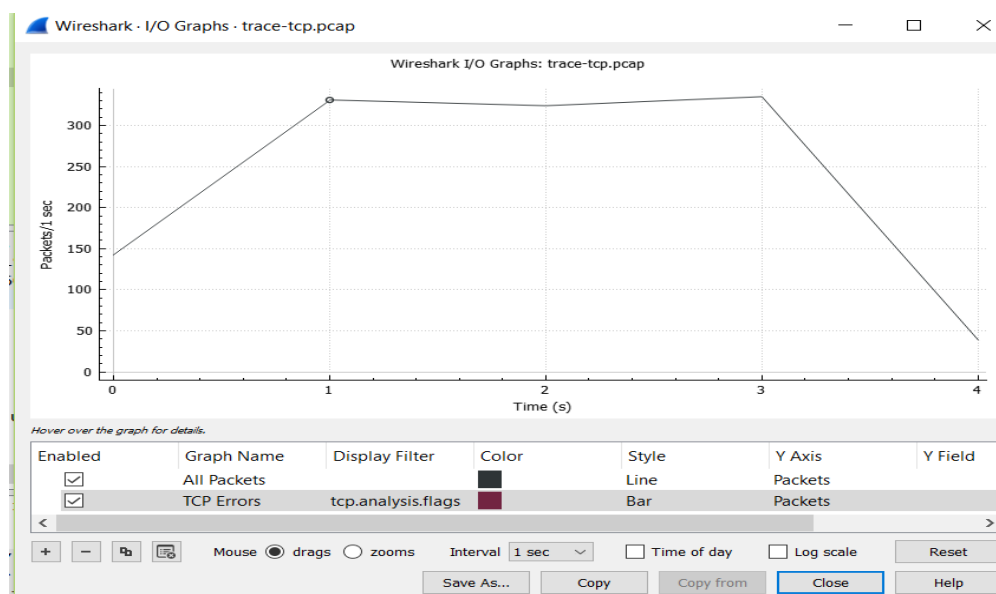


Figure 12: The IO graph

Now we will tweak it to show the download rate with the changes given below

- For the x-axis, adjust the interval to 100ms. The tick interval should be small enough to see into the behavior over the trace, and not so small that there is no averaging. 0.1 seconds is a good choice
- On the y-axis, change the unit to be Bits. The default is Packets. By changing it, we can easily work out the bits/sec throughput by taking the y-axis value and scaling as appropriate.
- Add a filter expression in the Display Filter to see only the download packets. So far we are looking at all of the packets. Assuming the download is from the usual web server port of 80, you can filter for it with a filter of `"tcp.srcport==80"`.
"Reset" button to cause it to redisplay.
- To see the corresponding graph for the upload traffic, press the plus button to enter a second filter in the next box. Again assuming the usual web server port, the filter is `"tcp.dstport==80"`.
- After you enable (checkbox) you should have two lines on the graph (Figure 13).

From Figure 13 below, we can see the sample down-load rate quickly increase from zero to a steady rate, with a bit of an exponential curve.

What is this exponential phase called in TCP? [Slow start](#)

What is the observed download rate when the connection is running? [2500000 bits/s = 312.5 KB](#)

Our download proceeds fairly steadily until it is done. This is the ideal, but many downloads may display more variable behavior if, for example, the available bandwidth varies due to competing downloads, the download rate is set by the server rather than the network, or enough packets are lost to disrupt the transfer.

On the other hand, the upload rate is also steady, and looks like a small trickle of ACK traffic. For e.g., the data rate in the upload direction in packets/second and bits/second due to the ACK packets is 120 packets/sec and 60,000 bits/sec, respectively.

Note, you can click on the graph to be taken to the nearest point in the trace if there is a feature you would like to investigate.

Try clicking on parts of the graph and watch where you are taken in the Wireshark trace window. 10



Figure 13: TCP download rate over time via an IO graph

Inspect the packets in the download in the middle of your trace for these features:

- You should see a pattern of TCP segments received carrying data and ACKs sent back to the server. Typically, there will be one ACK every couple of packets. These ACKs are called Delayed ACKs. By delaying for a short while, the number of ACKs is halved.
- Since this is a download, the sequence number of received segments will increase; the ACK number of subsequently transmitted segments will increase correspondingly.
- Since this is a download, the sequence number of transmitted segments will not increase (after the initial get). Thus the ACK number on received segments will not increase either.
- Each segment carries Window information to tell the other end how much space remains in the buffer. The Window must be greater than zero, or the connection will be stalled by flow control.

Typically the Ack number tells the next expected sequence number therefore it will be the sequence number of the next byte the receiver expects to receive.