

ICT2203/ICT2203X

Network Security



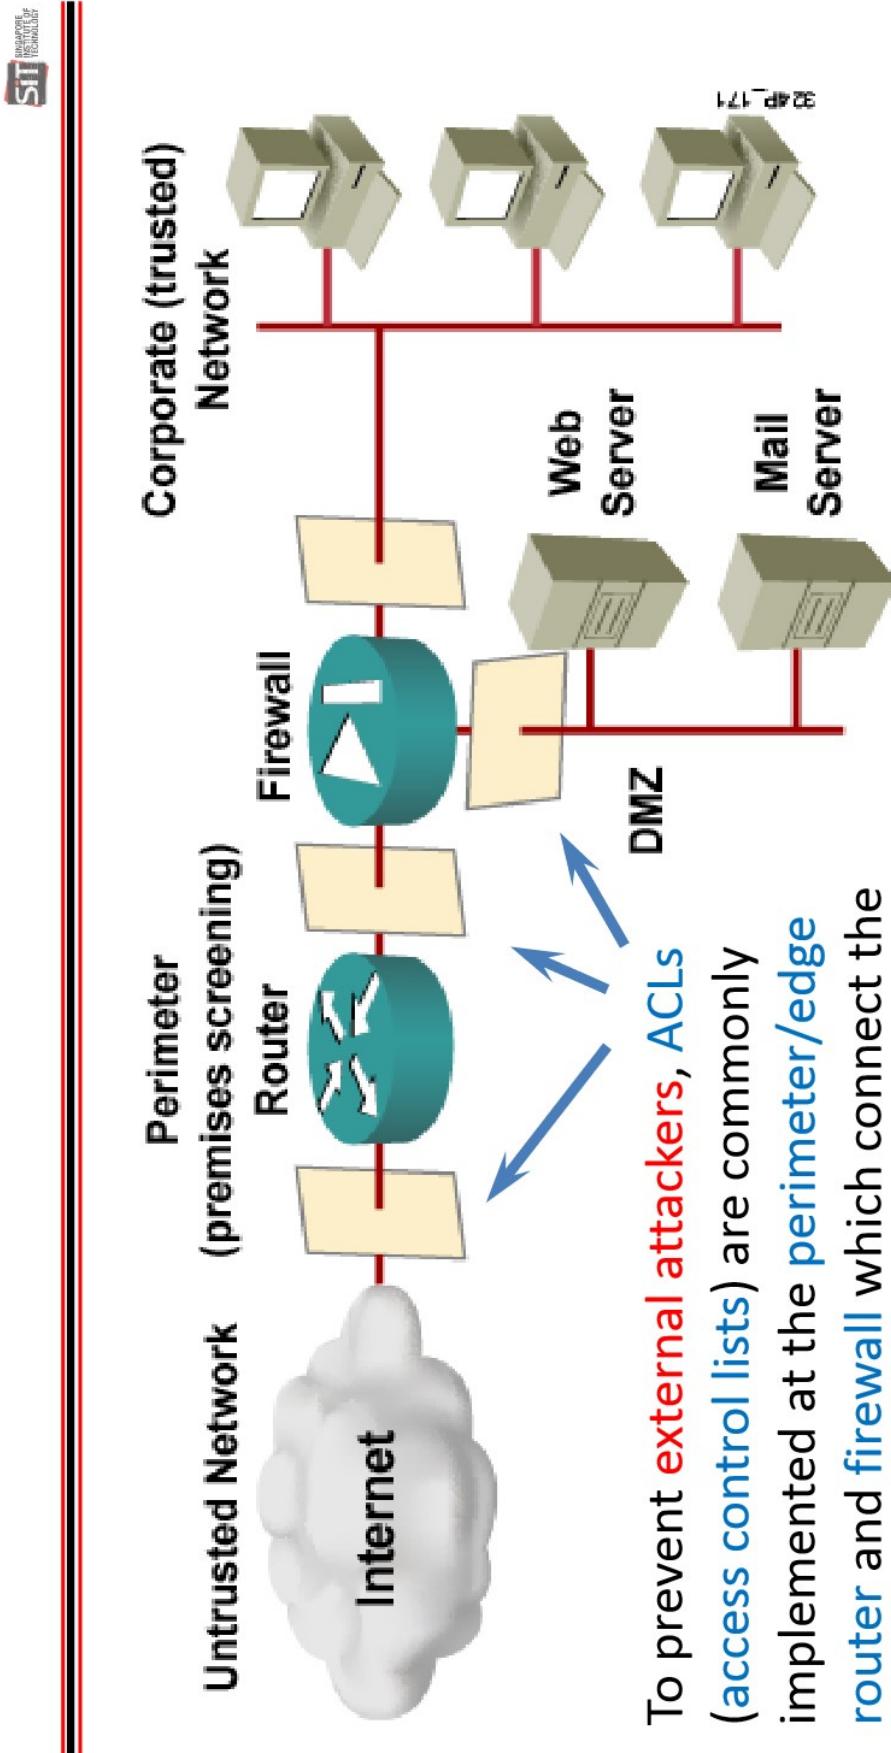
Lab 5 Notes:

Attacks and Defense of IP Networks with Routers

2021-2022 Trimester 3



In contrast to internal attacks, the first point of entry for **external attackers** are typically the **routers** or **firewalls**.



Lab Exercise 1



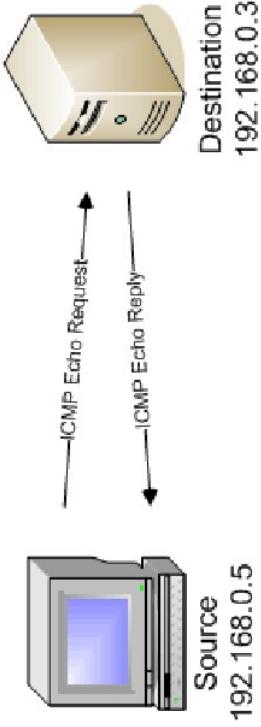
- 1.1 Network probing and scanning attack**

- 1.2 Defense:**
 - Minimise using suitable **ACLs** in routers

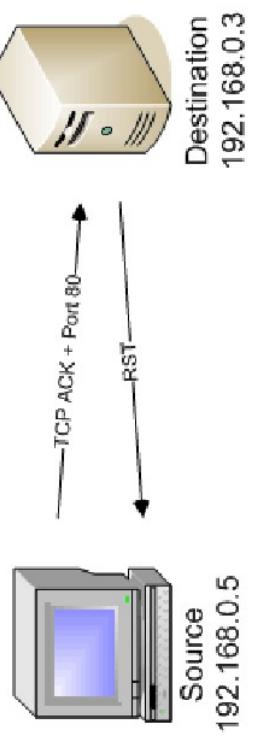
A potential external attack is **reconnaissance attack**, specifically **network probing** and **port scanning attacks** to discover targets which you'll learn more in ICT2204 Ethical Hacking.

Technically, **network probing** and **scanning attacks** are based on sending **ARP**, **ICMP**, **TCP** or **UDP** requests and receiving/not receiving the replies.

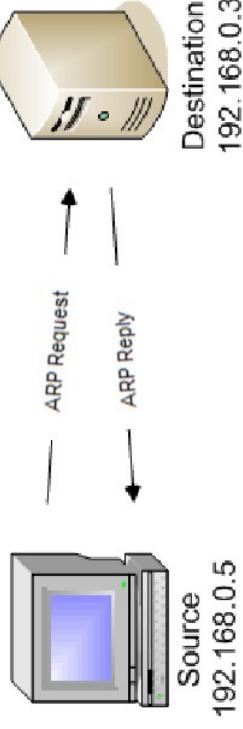
Example of **ICMP scan**:



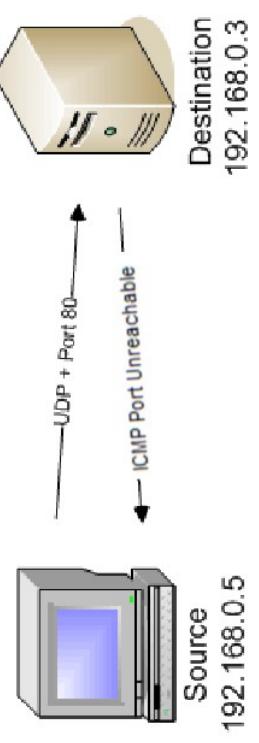
Example of **TCP scan**:



Example of **ARP scan**:



Example of **UDP scan**:



A useful tool for conducting network probing and scanning attacks is nmap which you will try in the lab.

Nmap - "Network Mapper"

Base Syntax

```
# nmap [ScanType] [Options] {targets}
```

Target Specification

IPv4 address: 192.168.1.1
IPv6 address: AABB:CCDD::FF%eth0
Host name: www.target.tgt
IP address range: 192.168.0.0-255
CIDR block: 192.168.0.0/16
Use file with lists of targets: -iL <filename>

Probing Options

- Pn Don't probe (assume all hosts are up)
- PB Default probe (TCP 80, 443 & ICMP)
- PS<portlist>
- PE Use ICMP Echo Request
- PP Use ICMP Timestamp Request
- PM Use ICMP Netmask Request

Target Ports

No port range specified scans 7,000 most popular ports

- F Scan 100 most popular ports
- P<port1>-<port2> Port range
- P<port1>,<port2>,... Port list
- PU:110,T20-445 Mix TCP and UDP
- r Scan linearly (do not randomize ports)
- top-ports <n> Scan in most popular ports
- p-65535 Leaving off initial port makes Nmap scan start at port 1
- pO- Leaving off end port makes Nmap scan up to port 65535
- p- Leaving off start and end port makes Nmap scan ports 1-65535

Aggregate Timing Options

- T0 Paranoid: Very slow, used for IDS evasion
- T1 Sneaky: Quite slow, used for IDS evasion
- T2 Polite: Slows down to consume less bandwidth, runs ~10 times slower than default
- T3 Normal: Default, a dynamic timing model based on target responsiveness
- T4 Aggressive: Assumes a fast and reliable network and may overwhelm targets
- T5 Insane: Very aggressive; will likely overwhelm targets or miss open ports

Output Formats

- ON Standard Nmap output
- OG Greppable format
- OX XML format
- OA <basename>

Scan Types

- Ss Probe only (host discovery, not port scan)
- ST SYN Scan
- TU UDP Scan
- SV Version Scan
- O OS Detection
- scanflags Set custom list of TCP using URGACKPSHRSSTSYNFIN in any order

Misc Options

- n Disable reverse IP address lookups
- b Use IPv6 only
- A Use several features, including OS Detection, Version Detection, Script Scanning (default), and traceroute
- reason Display reason Nmap thinks port is open, closed, or filtered



Aggregating Timing Options

- T0 Paranoid: Very slow, used for IDS evasion
- T1 Sneaky: Quite slow, used for IDS evasion
- T2 Polite: Slows down to consume less bandwidth, runs ~10 times slower than default
- T3 Normal: Default, a dynamic timing model based on target responsiveness
- T4 Aggressive: Assumes a fast and reliable network and may overwhelm targets
- T5 Insane: Very aggressive; will likely overwhelm targets or miss open ports

Output Formats

- ON Standard Nmap output
- OG Greppable format
- OX XML format
- OA <basename>

Scan Types

- Ss Probe only (host discovery, not port scan)
- ST SYN Scan
- TU UDP Scan
- SV Version Scan
- O OS Detection
- scanflags Set custom list of TCP using URGACKPSHRSSTSYNFIN in any order

Misc Options

- n Disable reverse IP address lookups
- b Use IPv6 only
- A Use several features, including OS Detection, Version Detection, Script Scanning (default), and traceroute
- reason Display reason Nmap thinks port is open, closed, or filtered

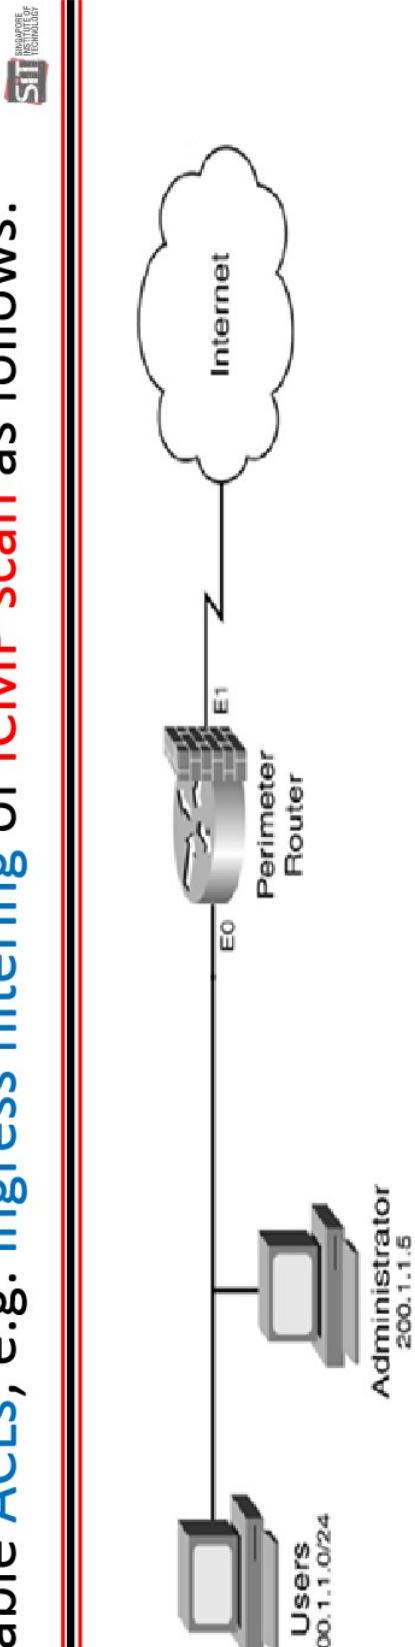
SANS www.sans.org

NMAP SECURITY 500  **PENETRATION TESTING CURRICULUM**
Free Resources: Blogs, Posters, Cheat Sheets pen-testing.sans.org 



5

Unfortunately, it can be difficult to fully prevent **network probing** and **scanning attacks**; but some may be blocked by suitable **ACLs**; e.g. ingress filtering of **ICMP scan** as follows:

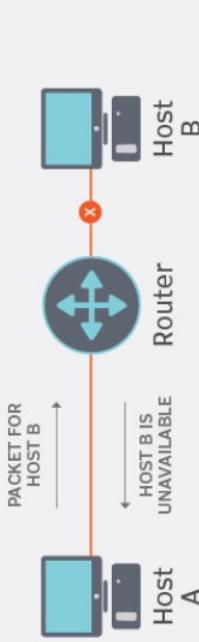


```
R1 (config) # ip access-list extended nmapICMPscan-in
R1 (config-ext-nacl) # deny icmp any any echo
R1 (config-ext-nacl) # deny icmp any any timestamp-request
R1 (config-ext-nacl) # deny icmp any any mask-request
R1 (config-ext-nacl) # permit ip any any
R1 (config) # interface E1
R1 (config-if) # ip access-group nmapICMPscan-in in
R1 (config-if) # no ip unreachables
```

Recall ICT1010 on ICMP error-reporting messages? Note that there are different **type 3 destination unreachable** error messages for different error conditions (codes) as shown:

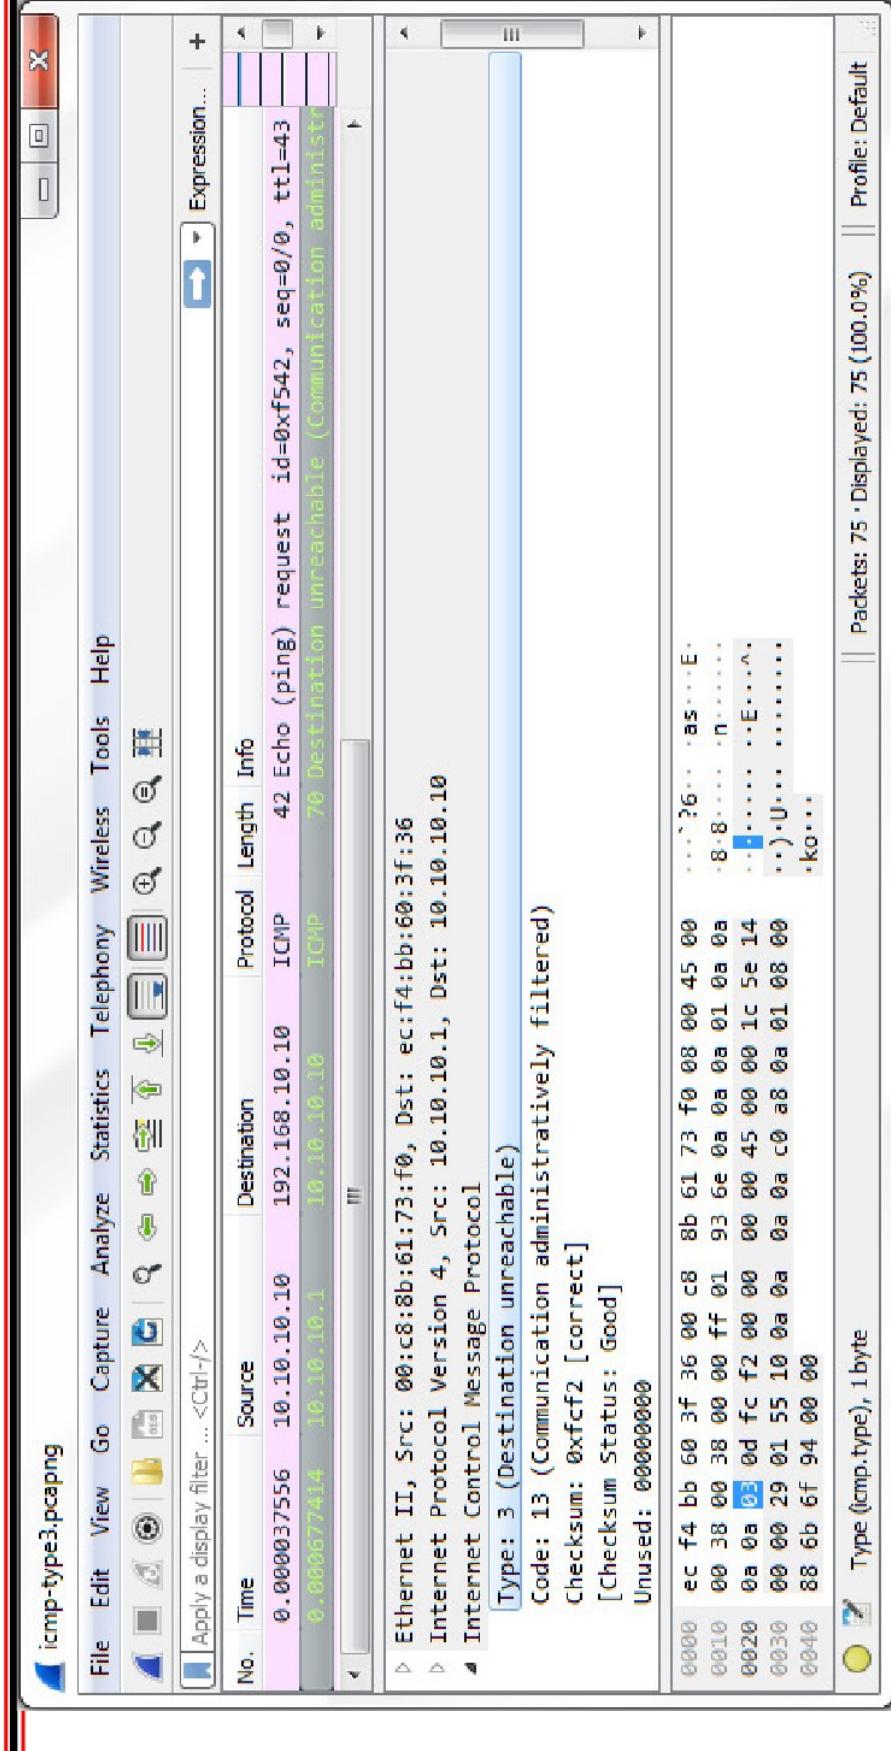
Type	Code	Description
3 Destination Unreachable	0	Destination network unreachable
	1	Destination host unreachable
	2	Destination protocol unreachable
	3	Destination port unreachable
	4	Fragmentation required, and DF flag set
	5	Source route failed
	6	Destination network unknown
	7	Destination host unknown
	8	Source host isolated
	9	Network administratively prohibited
	10	Host administratively prohibited
	11	Network unreachable for ToS
	12	Host unreachable for ToS
	13	Communication administratively prohibited
	14	Host Precedence Violation
	15	Precedence cutoff in effect

ICMP destination unreachable message

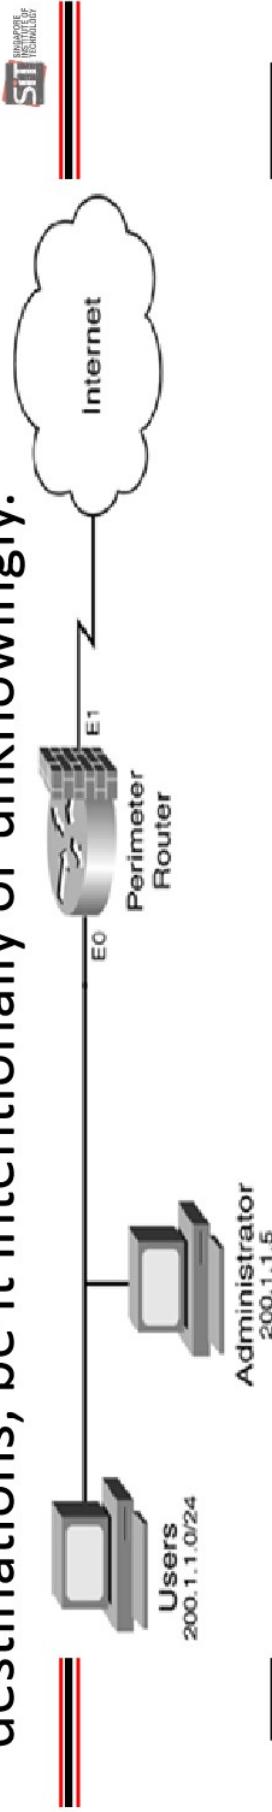


Host A attempts to send an IP packet to Host B. Host B is **unreachable**, so the router responds with a **Type 3 ICMP message**.

No ip unreachables will prevent a router configured with ACL to drop packets from sending ICMP type 3 destination unreachable which will leak information to attackers, e.g. shown below:



In addition, you may want to implement **egress filtering** of ICMP to prevent someone inside your network from **scanning** outside destinations, be it intentionally or unknowingly.



```
R1 (config) # ip access-list extended ICMPscan-out
R1 (config-ext-nacl) # permit icmp host 200.1.1.5 any echo ①
R1 (config-ext-nacl) # permit icmp 200.1.1.0 0.0.0.255 any
replace-with-actual-ICMP-error-reporting-message-see-below ②
R1 (config-ext-nacl) # deny icmp any any
R1 (config-ext-nacl) # permit ip any any
R1 (config) # interface E1
R1 (config-if) # ip access-group ICMPscan-out out
```

- ① Allow administrator to ping which is useful for testing connectivity.
- ② Trade-off allowing ICMP error-reporting messages vs **attacks**; e.g.
 - **unreachable**: inform destination unreachable but exploited by **UDP scan**
 - **time-exceeded**: exploited by attacker to **trace route** to target

Lab Exercise 2



2.1

Smurf attack with **spoofed source IP address**

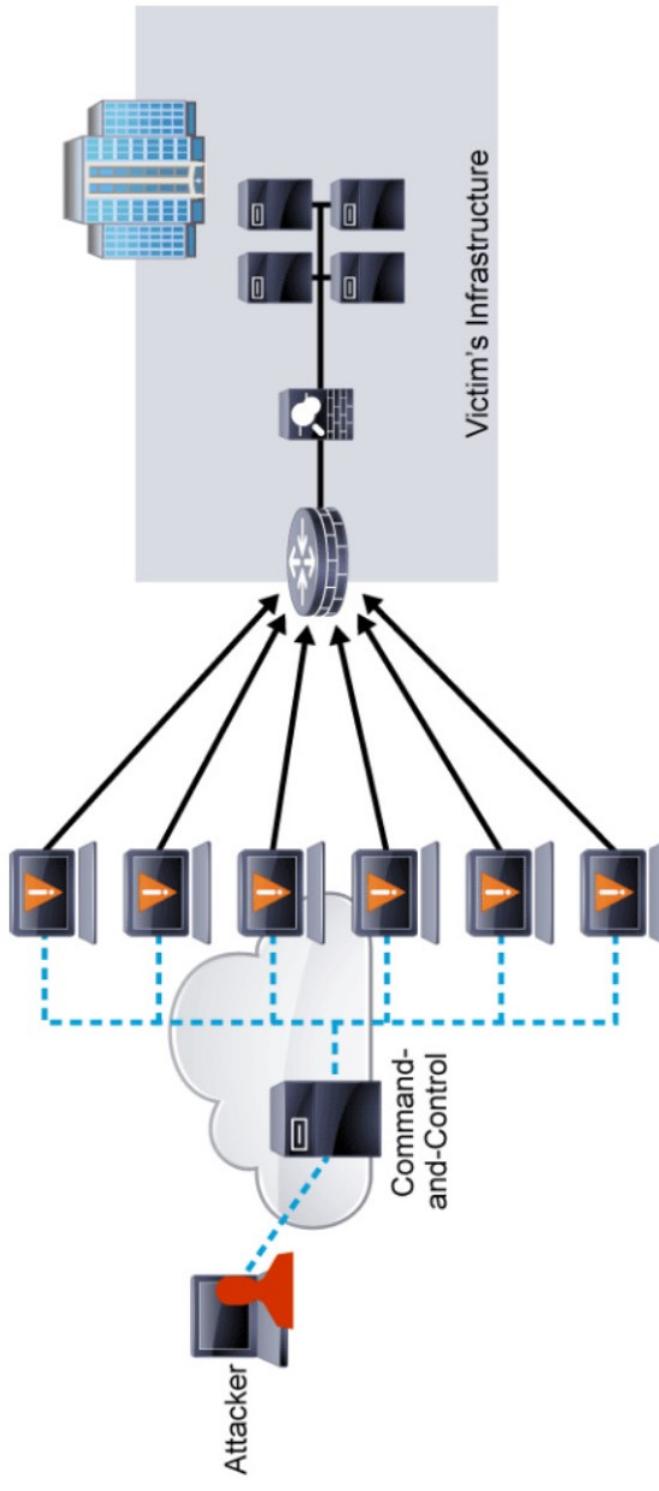
2.2

Defenses:

- Disable hosts from replying to **broadcast ping**
- Disable routers from supporting **directed broadcast**
- Apply **BCP 38/RFC 2827** and **anti-spoofing ACLs** in routers

Another potential external attack is **DoS (Denial-of-Service)** or **DDoS (Distributed DoS)** which aims to make the target unavailable to authorised users.

Typically, **DoS/DDoS** aims to overload the **network bandwidth** or **resources** of the target to make it unavailable.



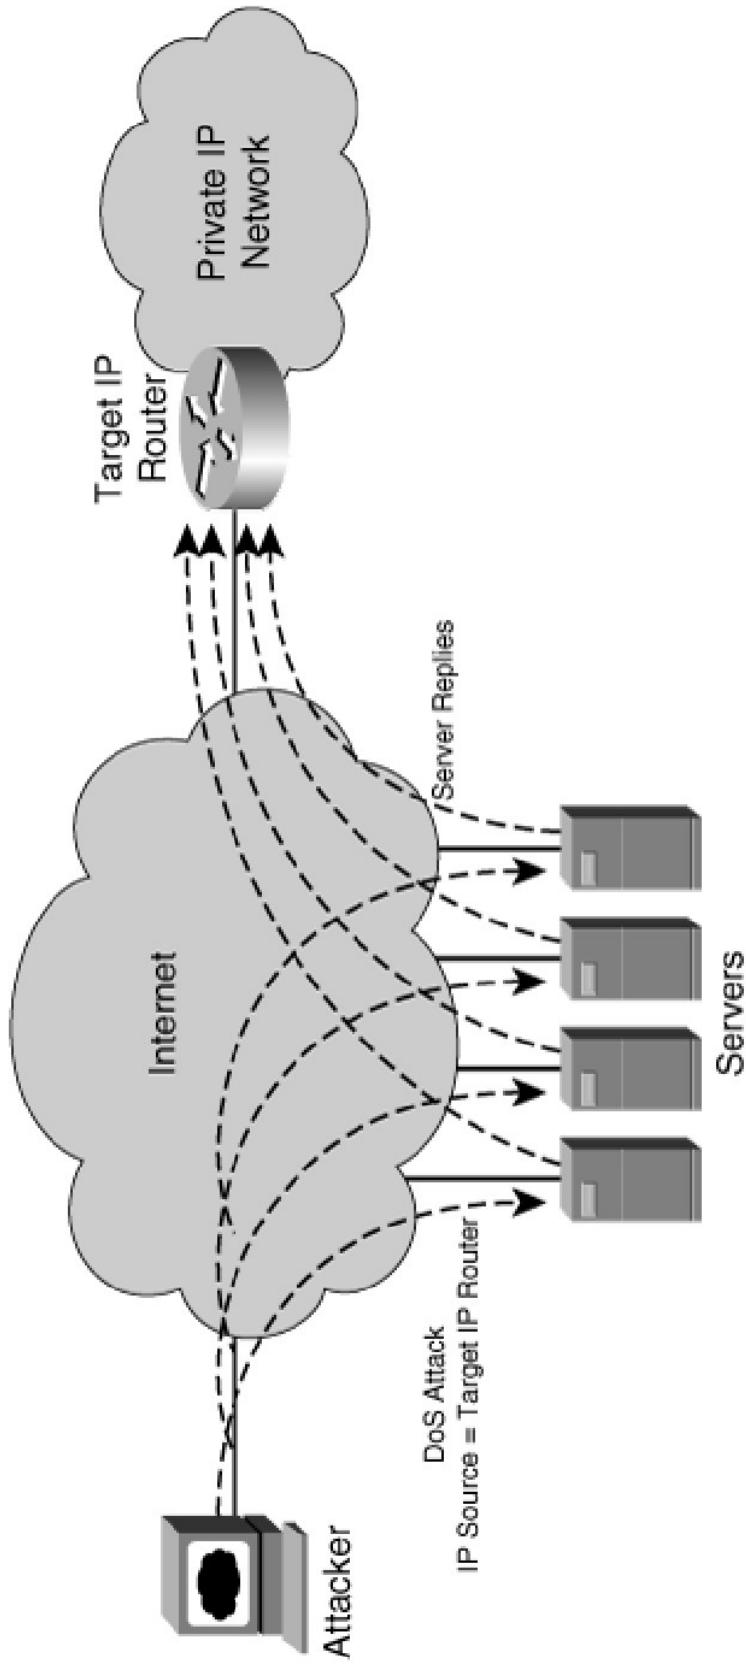
DDoS attack is basically **DoS** attacks that are launched using an army of compromised hosts called zombies/bots distributed in the Internet.

Broadly, **DoS/DDoS attacks** may be classified into 2 categories as follows:



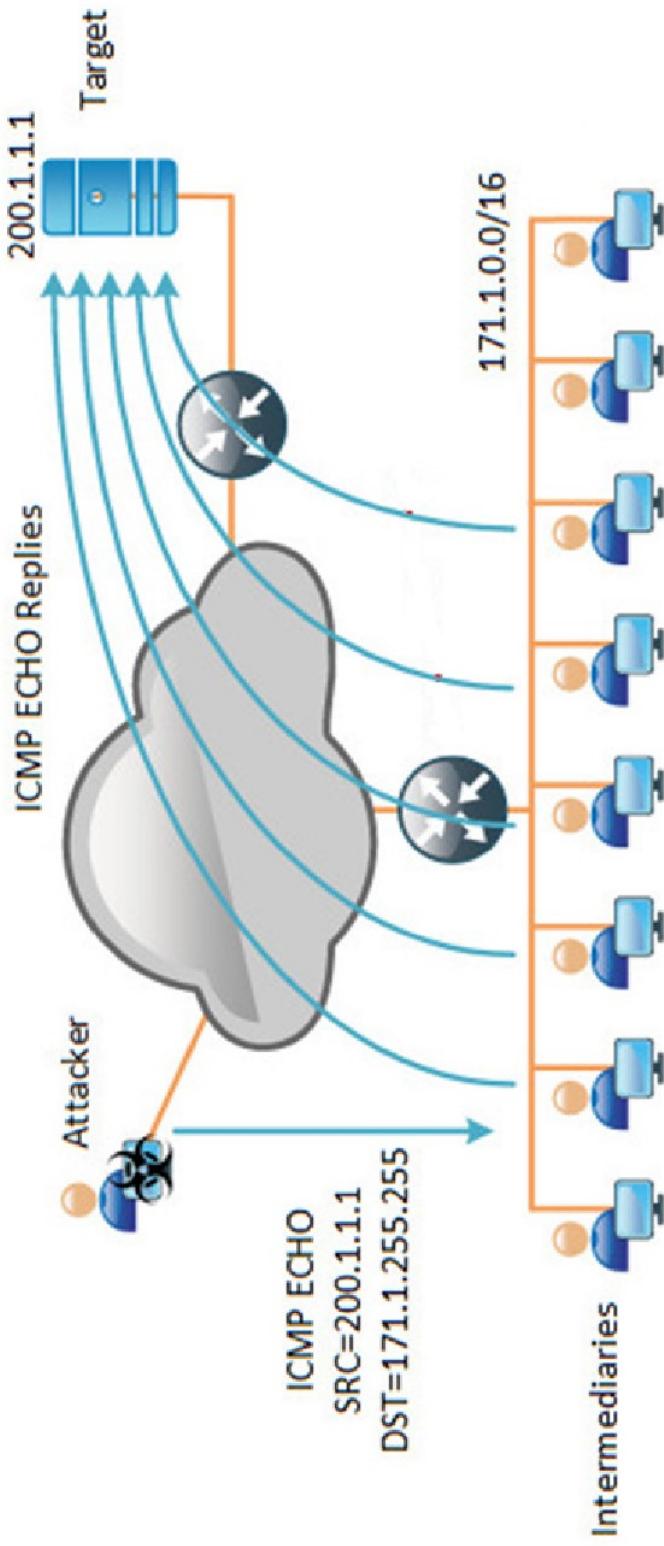
DoS/DDoS categories	Description	Typical attack packet rate	Examples
Overwhelming Quantity of Traffic	Threat actor sends an enormous quantity of data at a rate that the network, host, or application cannot handle, possibly causing slow response times or even crash a device or service.	High	ICMP flood, UDP flood, smurf attack, TCP SYN, etc
Maliciously Crafted Packets	Threat actor sends maliciously crafted packets that the host or application is not expected to handle, possibly causing it to run very slowly or even crash.	Low	Slowloris, ping of death, etc

Reflection attack is a type of DoS/DDoS which works by sending packets to intermediary servers using **spoofed source IP** belonging to target so that replies are **reflected** to attack target



If the reflection attack uses small forged packets which resulted in large replies from the intermediaries, it is also known as **amplification attack**; e.g. **smurf attack** appeared in late 1990s.

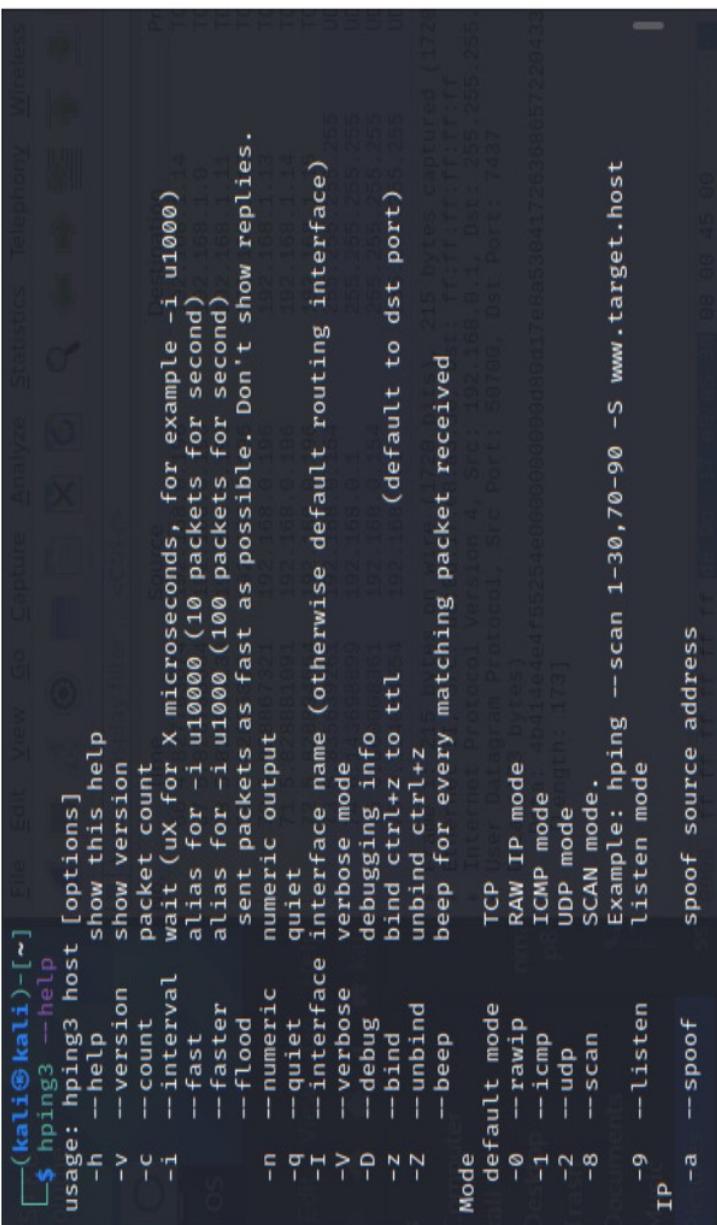
In **smurf attack**, a single **broadcast ping** sent to a **directed broadcast address** by attacker using **spoofed IP source address** belonging to target would result in multiple ping replies (each from an intermediary in the directed broadcast subnet) reflected to the target.



To have a better understanding of **amplification attack**, you will use **hping3** in Kali Linux to try **smurf attack** in the lab.

```
root@kali:~# hping3 -1 --flood -a target_ip directed_broadcast_address
```

Refer <https://www.kali.org/tools/hping3> on the use of hping3:



```
(kali㉿kali)-[~] $ hping3 --help
usage: hping3 host [options]
      -h --help          show this help
      -v --version        show version
      -c --count          packet count
      --interval         wait (ux for x microseconds, for example -i u1000)
      -i --fast           alias for -i u10000 (10 packets for second)
      --faster           alias for -i u1000 (100 packets for second)
      --flood            sent packets as fast as possible. Don't show replies.
      -n --numeric        numeric output
      -q --quiet          quiet
      -I --interface      interface name (otherwise default routing interface)
      -V --verbose         verbose mode
      -D --debug           debugging info
      -z --bind            bind ctrl+z to ttl
      -Z --unbind          unbind ctrl+z
      --beep              beep for every matching packet received
Mode
default mode
      -0 --rawip          TCP User Datagram Protocol, Src: 192.168.0.1, Dst: 255.255.255.255
      -1 --icmp           RAW IP mode (bytes)
      -2 --udp             ICMP mode (ab1484e4f55254e0000000000d80cd17e8a530e4172636865722943)
      -8 --scan            UDP mode
      -9 --listen          SCAN mode.
Example: hping --scan 1-30,70-90 -S www.target.host
      IP                spoof source address
-a --spoof
```

Today, **smurf attack** has been largely prevented by default through:

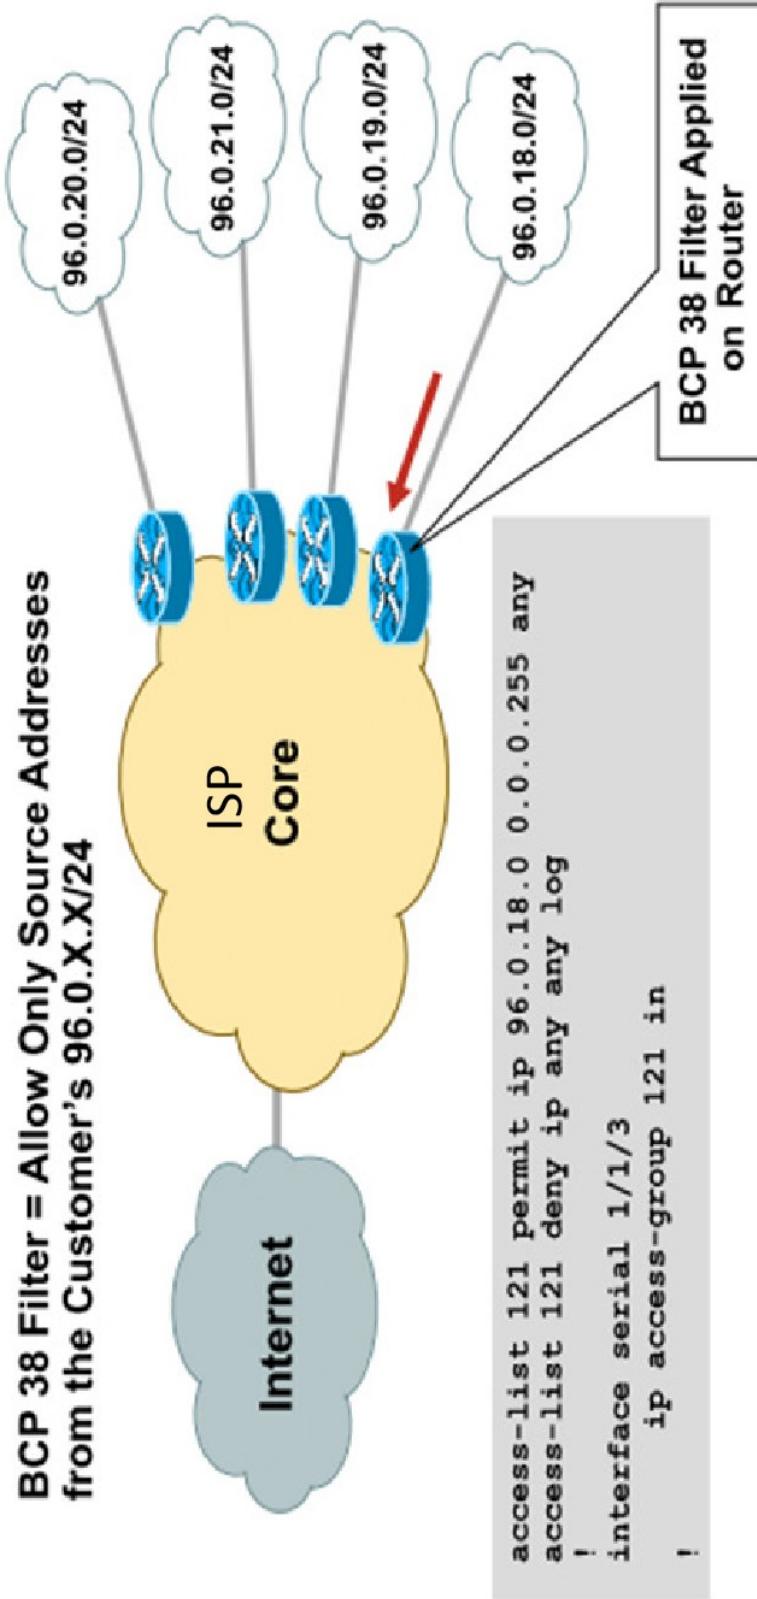
- Disable routers from supporting **directed broadcast**;
- Disable hosts from replying to **broadcast ping**.

To prevent DoS/DDoS attack using spoofed source IP addresses

like smurf attack, ISPs today are recommended to implement

Best Current Practice BCP 38/RFC 2827.

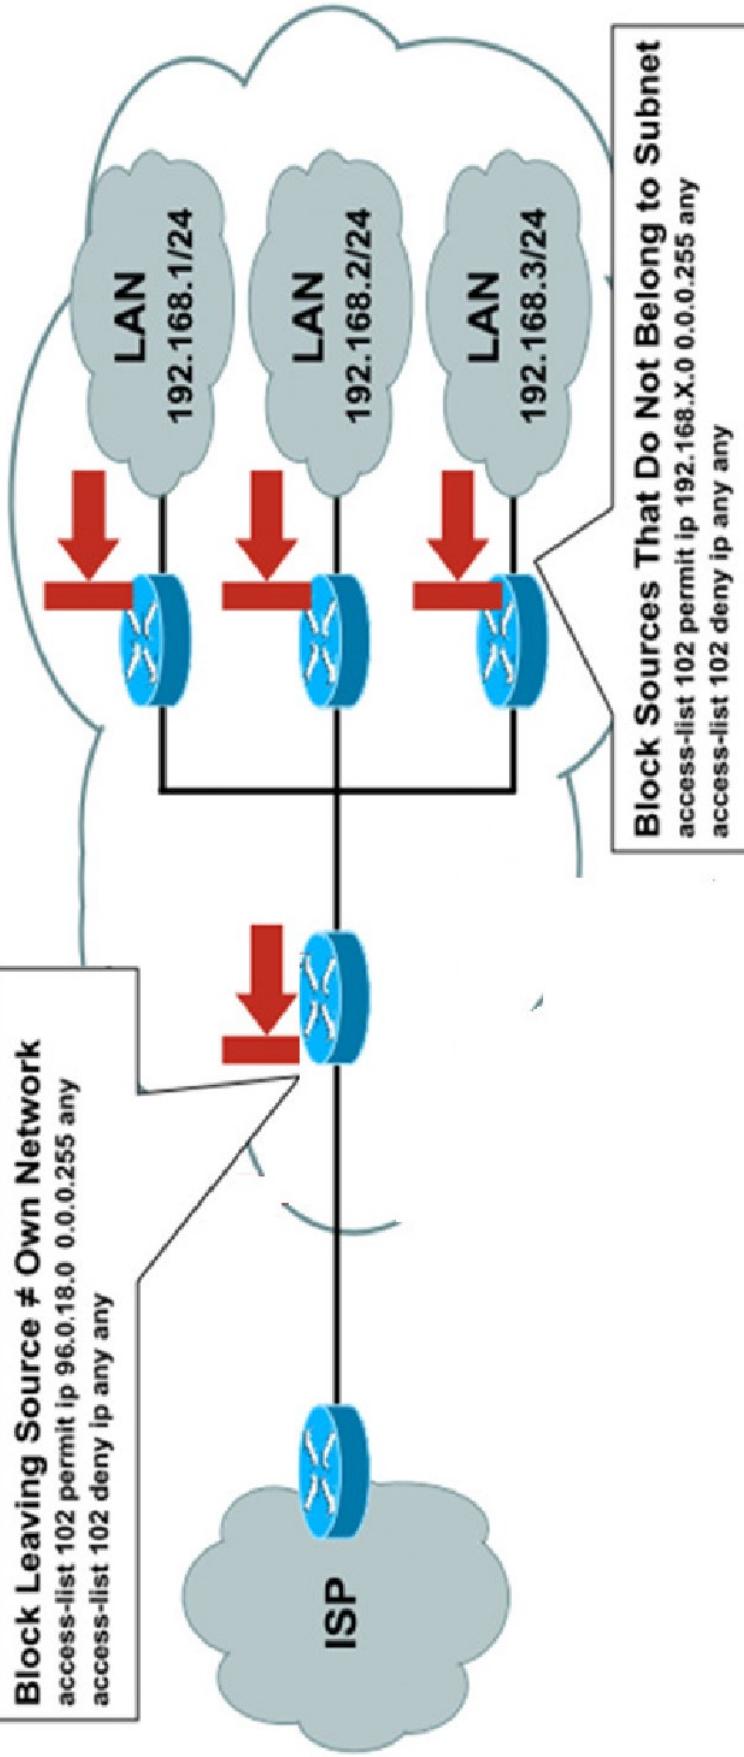
BCP 38/RFC 2827 recommends ISPs to implement **network ingress filtering** to prevent their networks from permitting spoofed IP addresses and supporting DoS/DDoS attacks.



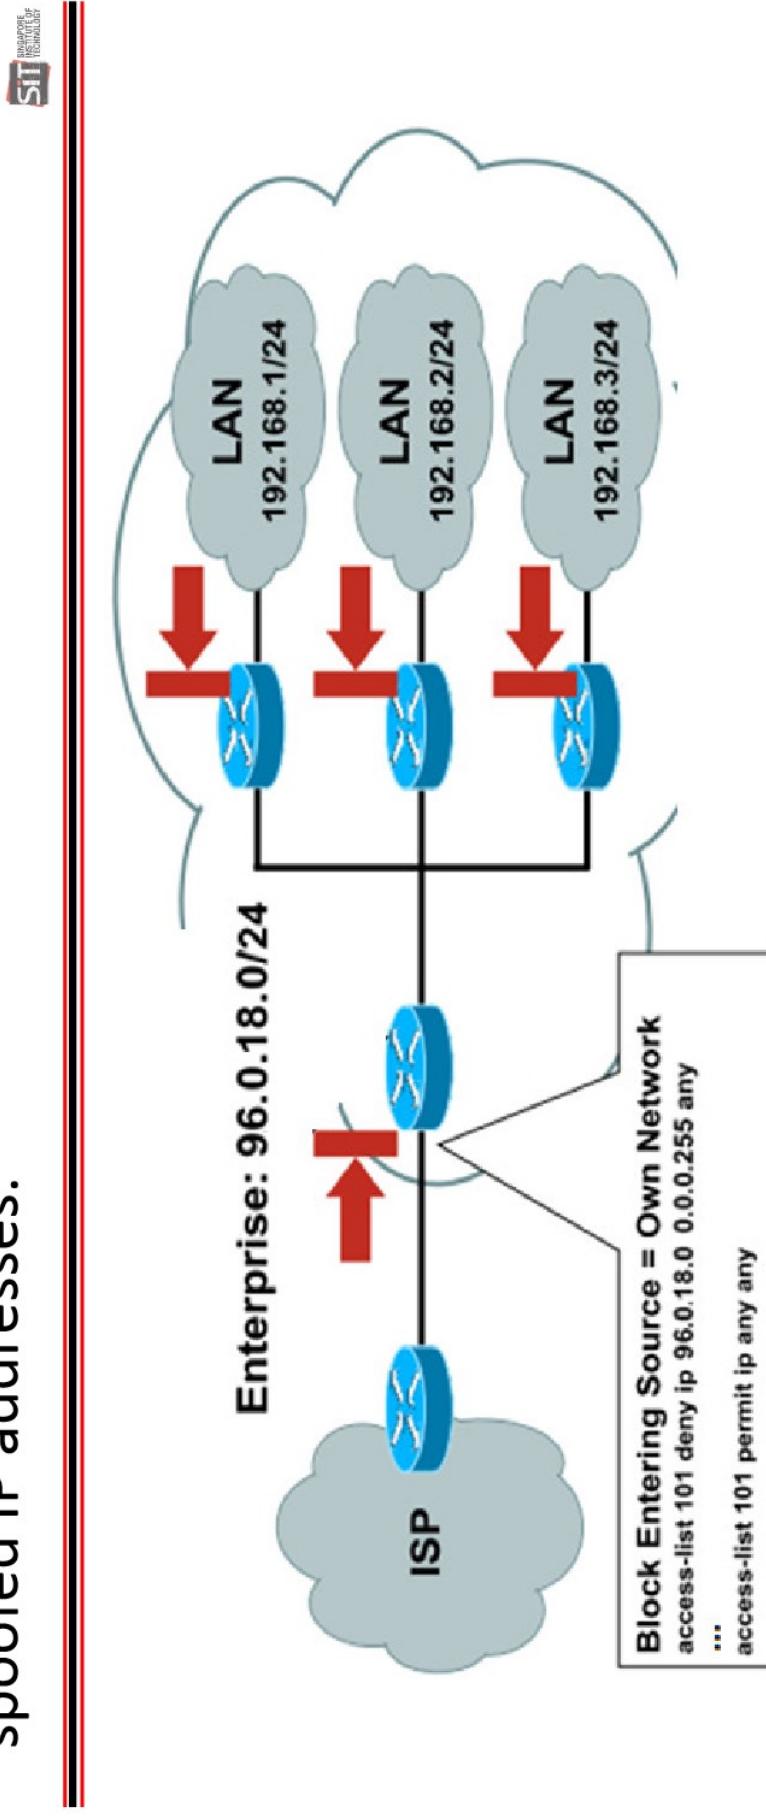
Similarly, organizations should also implement **BCP 38/RFC 2827** at their edge routers to prevent their hosts, possibly zombies, from using **spoofed source IP addresses** to launch attacks.

Enterprise: 96.0.18.0/24

Block Leaving Source ≠ Own Network
access-list 102 permit ip 96.0.18.0 0.0.0.255 any
access-list 102 deny ip any any



In addition, organizations should implement **anti-spoofing ACL** at their **edge routers** to prevent being victim to attacks using spoofed IP addresses.



Anti-spoofing ACL is used to deny inbound IP packets that are obviously spoofed; e.g. source addresses that match your allocated public IP addresses, or special addresses in [RFC 3330](#) and [RFC 1918](#) that should not be used.

Typically, anti-spoofing ACL is implemented as a part of overall infrastructure protection ACL at the **edge router** for ingress filtering of malicious traffic from entering the network.



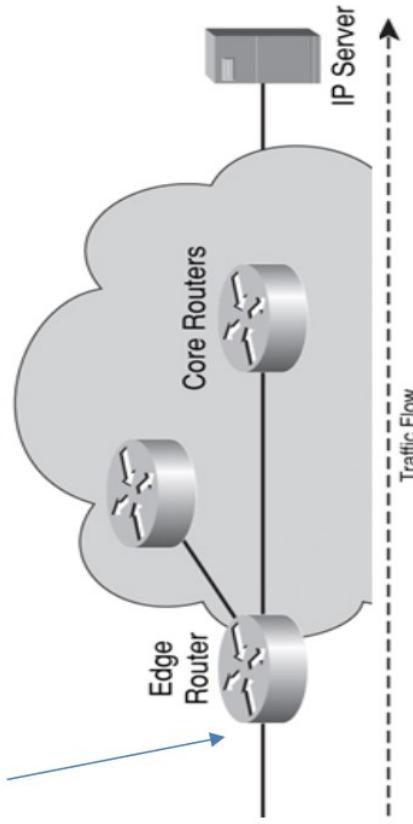
!— Anti-spoofing entries are shown here.

!— Deny RFC 3330 special-use address

```
access-list 110 deny ip host 0.0.0.0 any  
access-list 110 deny ip 127.0.0.0 0.255.255.255 any  
access-list 110 deny ip 192.0.2.0 0.0.0.255 any  
access-list 110 deny ip 224.0.0.0 31.255.255.255 any
```

!— Filter RFC 1918 space.

```
access-list 110 deny ip 10.0.0.0 0.255.255.255 any  
access-list 110 deny ip 172.16.0.0 0.15.255.255 any  
access-list 110 deny ip 192.168.0.0 0.0.255.255 any
```



(Infrastructure and anti-spoofing protection).
Applied on input of interface.

!— Deny your space as source from entering your AS. !— Deploy only at the AS edge.

```
access-list 110 deny ip YOUR_CIDR_BLOCK any
```

!— Deny access to internal infrastructure addresses.

```
access-list 110 deny ip any INTERNAL_INFRASTRUCTURE_ADDRESSES
```

!— Permit transit traffic.

```
access-list 110 permit ip any any
```

Lab Exercise 3



3.1

Tiny fragment and overlapping fragment attacks

3.2

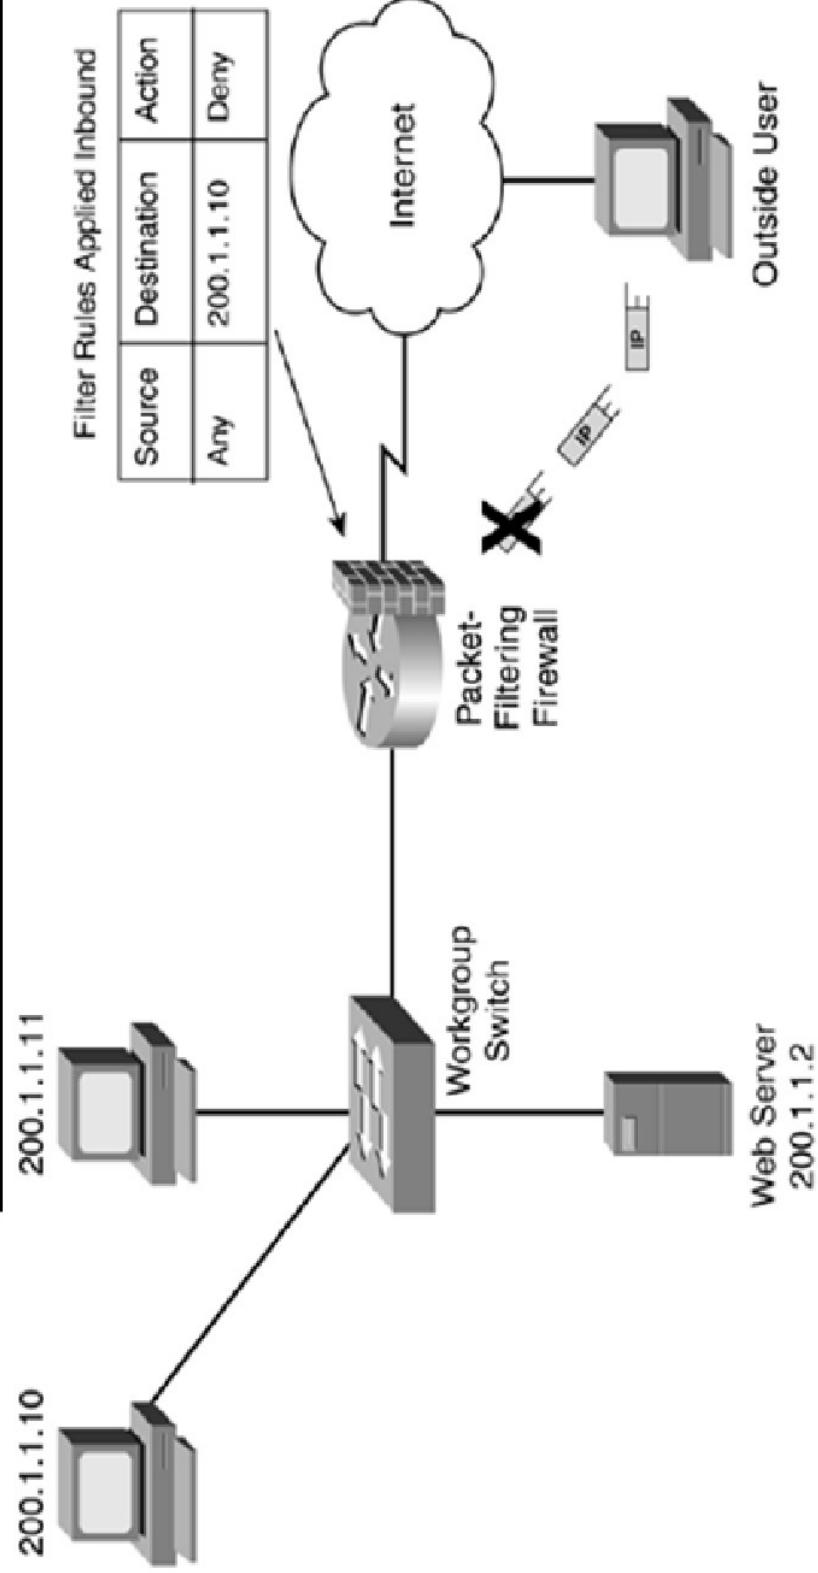
Defenses:

- RFC 1858 to discard fragments with offset $FO = 1$
- Apply ACLs with ‘fragments’ option

To protect internal **hosts**, one may be tempted to simply use **ACL** to filter external attackers from reaching them, e.g. for host 200.1.1.10 as follows:

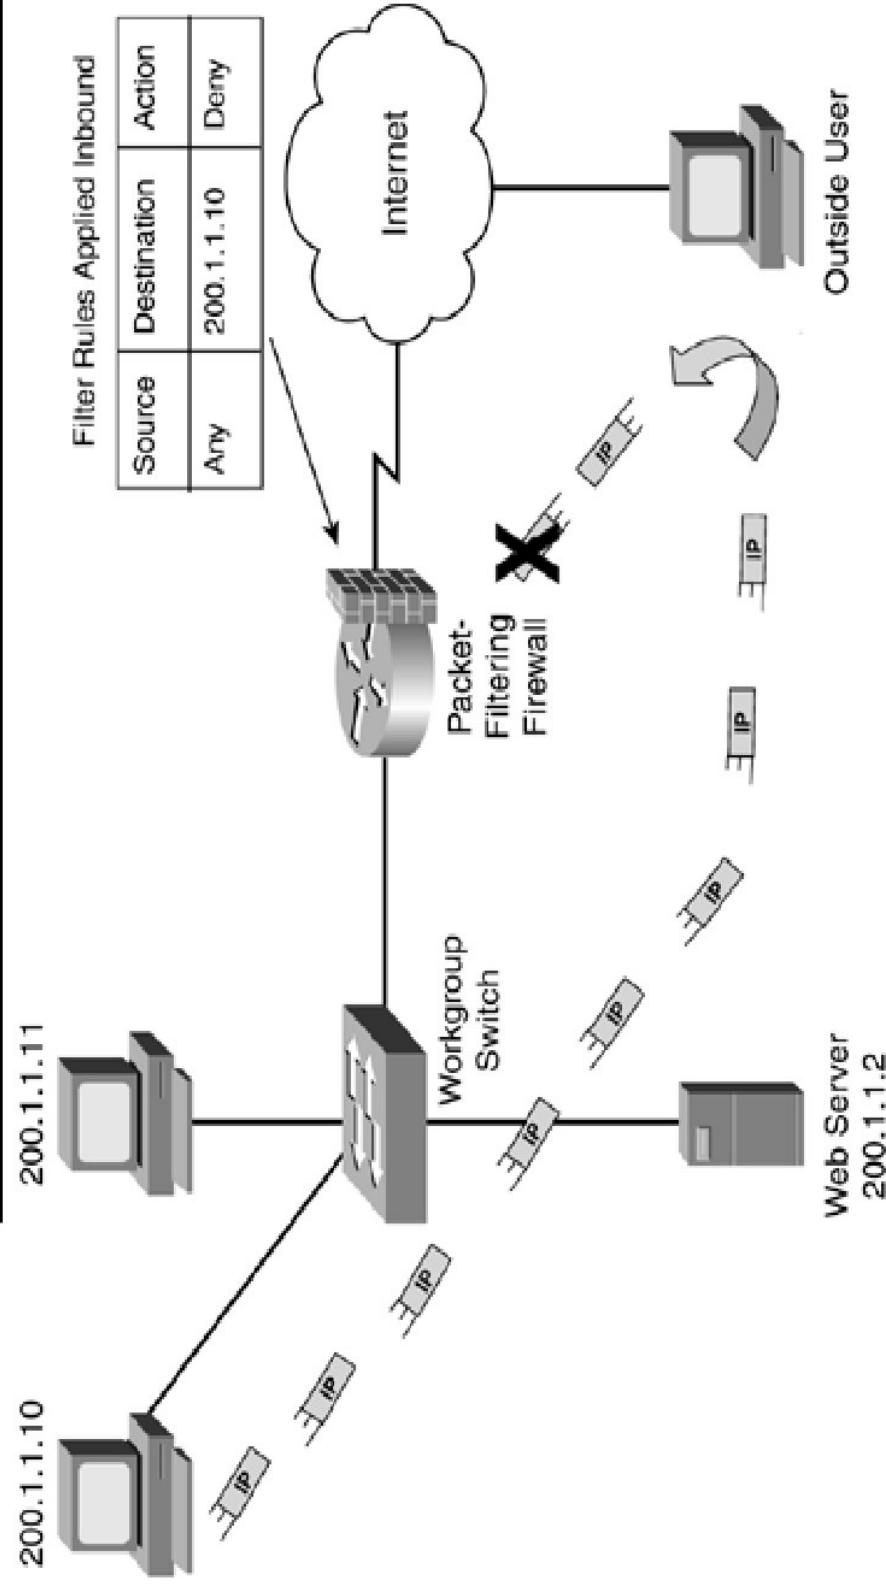


```
R1(config-ext-nacl)# deny tcp any host 200.1.1.10
```



However, **blocking external attackers** from reaching internal hosts will also lead to **blocking internal hosts** from accessing the Internet because all returning packets will be blocked!

```
R1(config-ext-nacl)# deny tcp any host 200.1.1.10
```

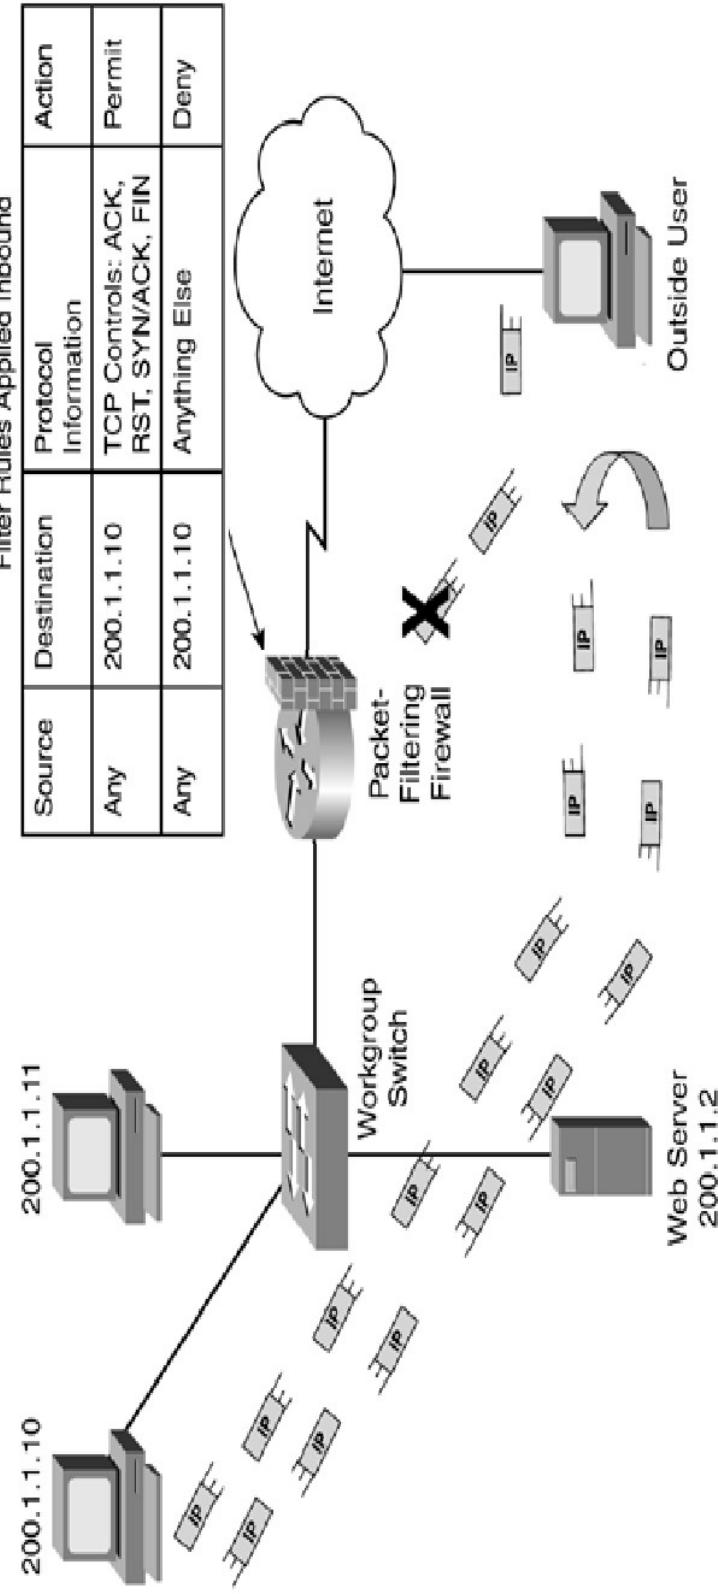


So, ACL statement has been enhanced with an optional '**‘established’**' parameter for protecting internal hosts.

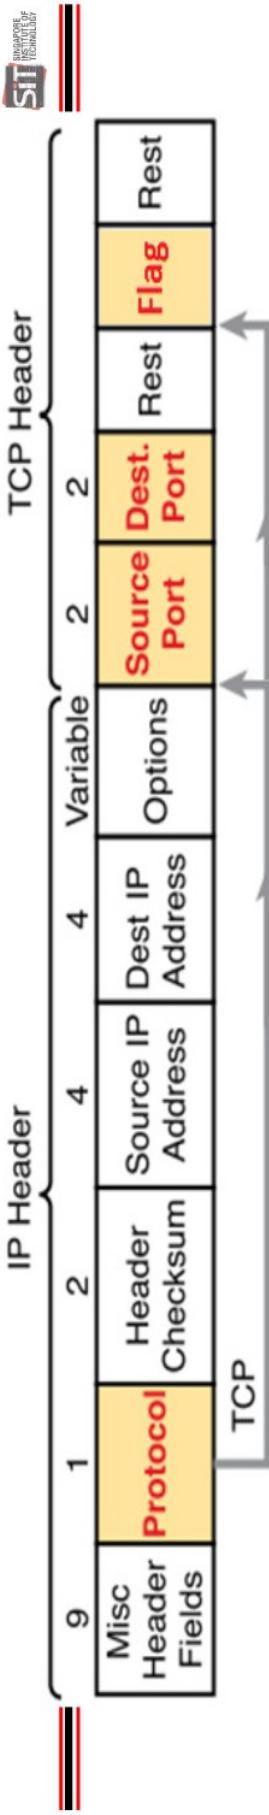
```
R1(config-ext-nacl)# permit tcp any host 200.1.10 established
```

The '**‘established’**' option applies only to **TCP** and is meant to:

- Allow returning TCP segments of **existing connections**
- **Block** TCP segments initiating a **new connection** from outside, or other malformed segments



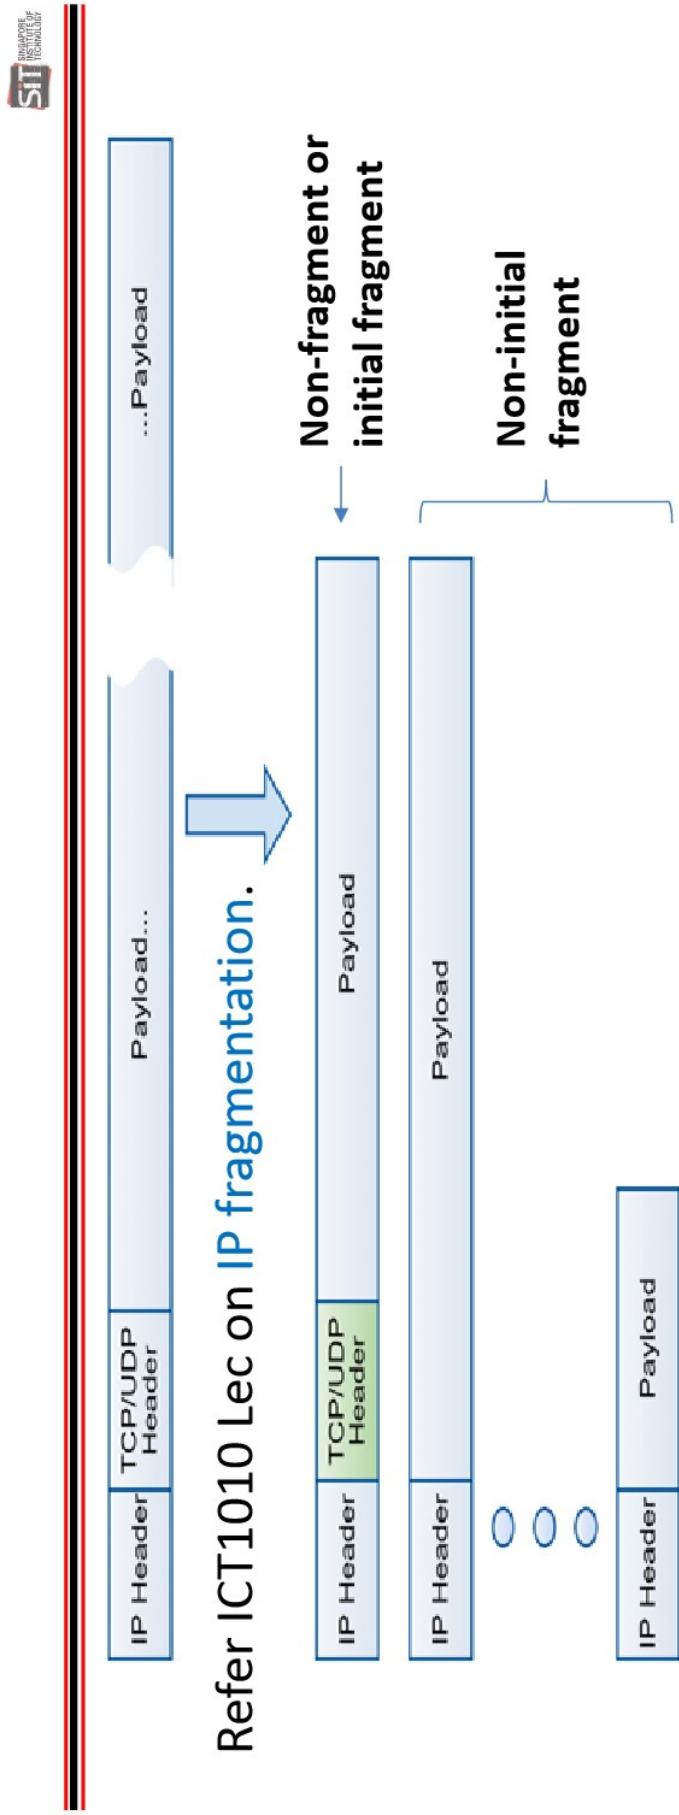
Alternatively, you may configure ACL to examine the **TCP flags** (refer ICT1010 Lec On TCP) of each packet, e.g. as follows:



- Allow returning TCP segments of **existing connections**, e.g. TCP flags with **+syn +ack, +ack, +fin, +fin +ack, +rst**
- Block** other TCP segments, e.g. **+syn** because it is the first segment of 3-way handshake initiating a **new connection**

```
R1 (config) # ip access-list extended TCPFLAG
R1 (config-ext-nacl) # permit tcp any any match-all +syn +ack
R1 (config-ext-nacl) # permit tcp any any match-all +ack
R1 (config-ext-nacl) # permit tcp any any match-all +fin
R1 (config-ext-nacl) # permit tcp any any match-all +fin +ack
R1 (config-ext-nacl) # permit tcp any any match-all +rst
R1 (config-ext-nacl) # deny tcp any any
R1 (config-ext-nacl) # permit ip any any
```

Unfortunately, attackers have come up with **IP fragmentation attacks** to **disguise IP packets** to **evasive ACL filtering**.



By default, **ACLs** are only applied to **non-fragments** and **initial fragments**.

For efficiency, **non-fragments** are simply passed without filtering because if **initial fragment** is already **filtered**, destination host will eventually discard non-initial fragments after reassembly timeout.

Specifically, one technique of using **IP fragmentation** to evade
ACL filtering is **tiny fragment attack** (RFC 1858).

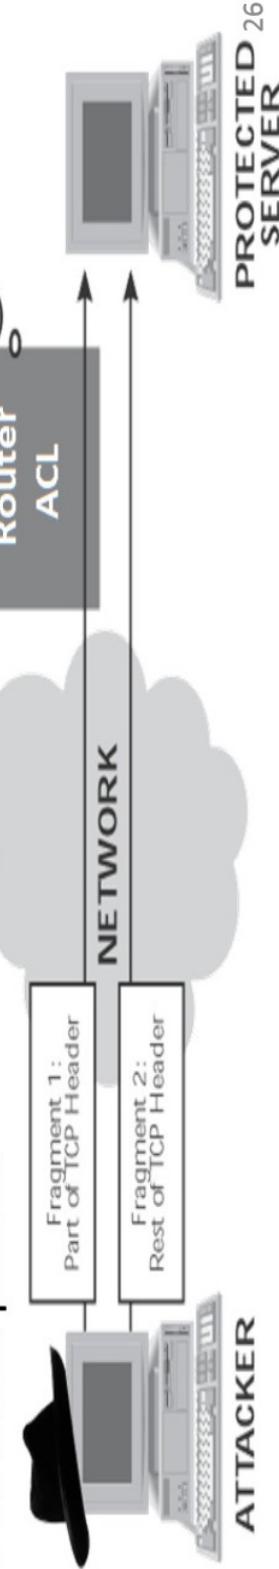
Tiny fragment attack intentionally breaks an IP packet into many **tiny
fragments**; e.g. 8 bytes of data:

Since TCP header is 20 bytes:

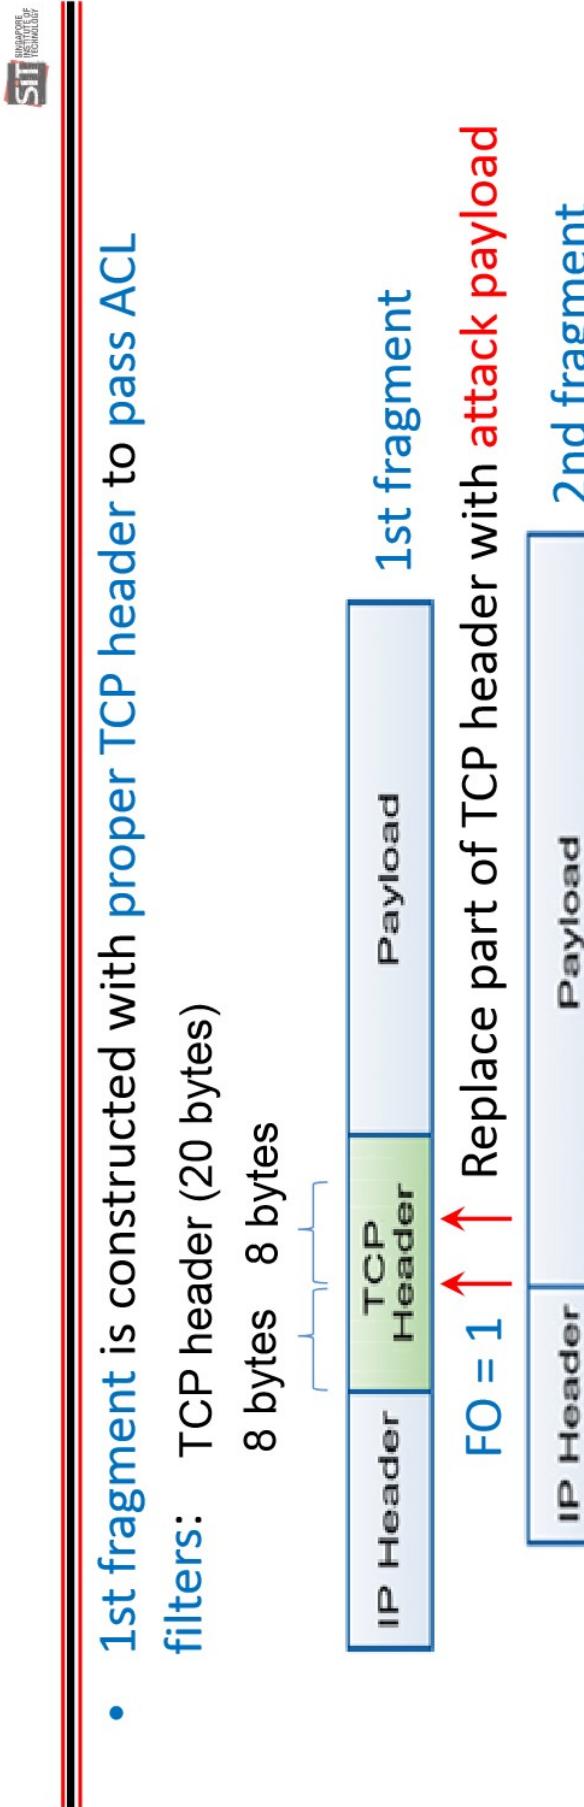


As a result:

- **1st fragment** will pass through because ACL that filters based on TCP flags will not be matched
- All **non-initial fragments** will also pass through
 - Thus all received and reassembled, which could be an attack packet!



Another technique of using IP fragmentation to **e evade ACL filtering** is **overlapping fragment attack** (RFC 1858).

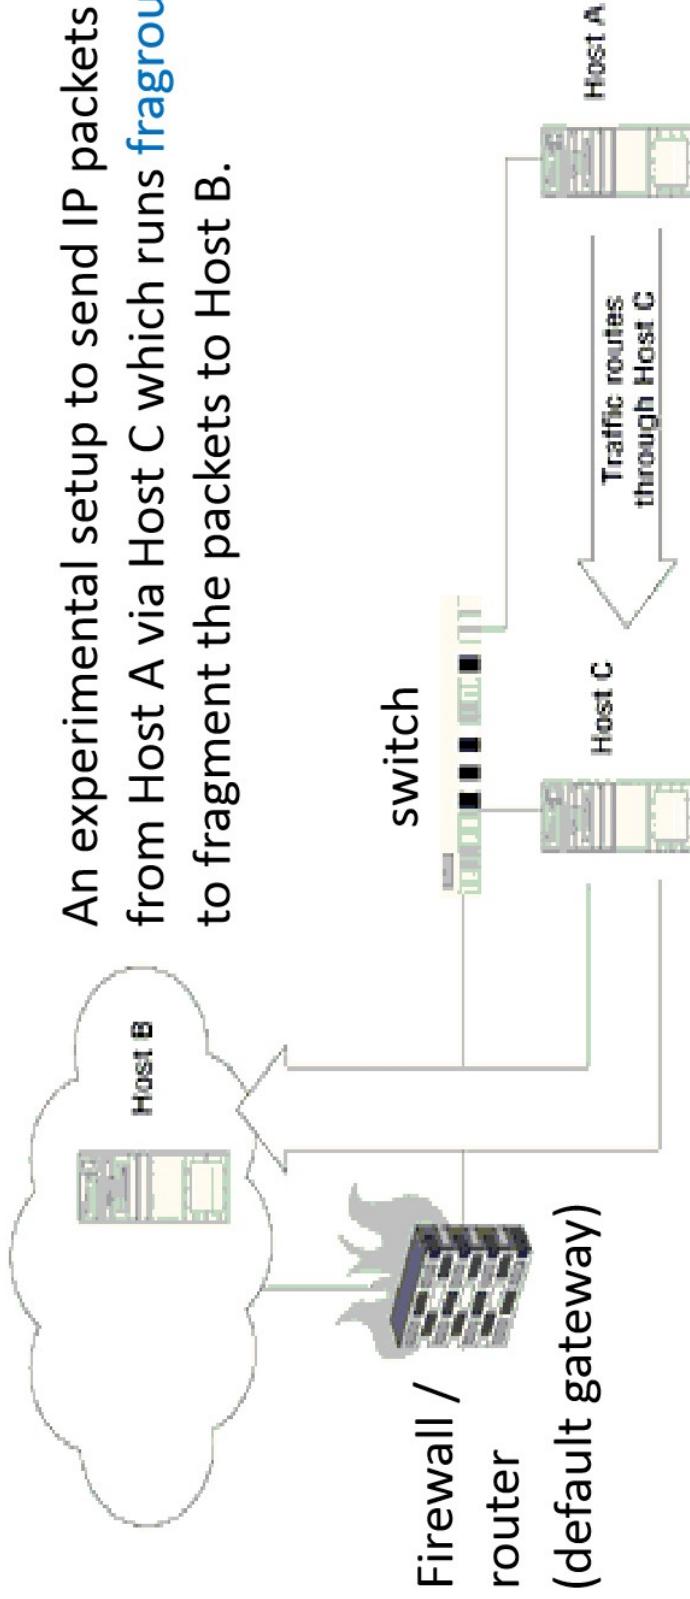


- 1st fragment is constructed with proper TCP header to pass ACL filters: TCP header (20 bytes)
8 bytes 8 bytes
- 2nd fragment is constructed with desired fragmentation offset (e.g. $FO = 1$) and **attack payload**
- When all received and reassembled, **attack payload** will overlap and replace part of original TCP header, leading to an attack packet!

To experiment with **IP fragmentation attacks**, you will try a tool called **fragrouter** in Kali Linux.

Refer <http://linux.die.net/man/8/fragrouter> on the use of **fragrouter**:

An experimental setup to send IP packets from Host A via Host C which runs **fragrouter** to fragment the packets to Host B.

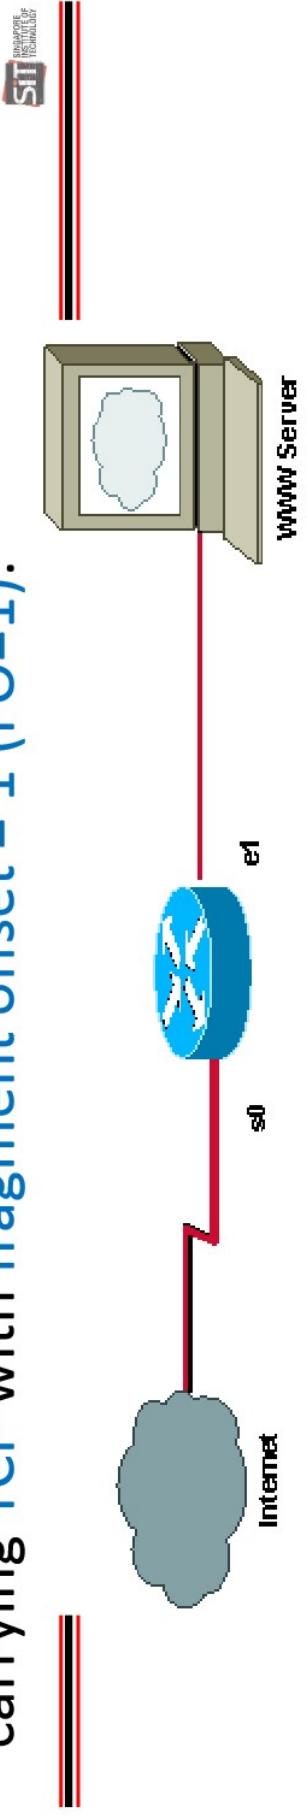


Host C default route to default gateway, running **fragrouter**, e.g. **tiny fragment attack**:

```
root@kali:~# fragrouter -i eth0 -F1
```

Change Host A default route to Host C IP address instead of default gateway.

To defend against **tiny fragment** and **overlapping fragment attacks**, Cisco has implemented RFC 1858 to drop IP fragment carrying TCP with fragment offset = 1 (FO=1).



Applying following ACL for ingress filtering at interface s0:

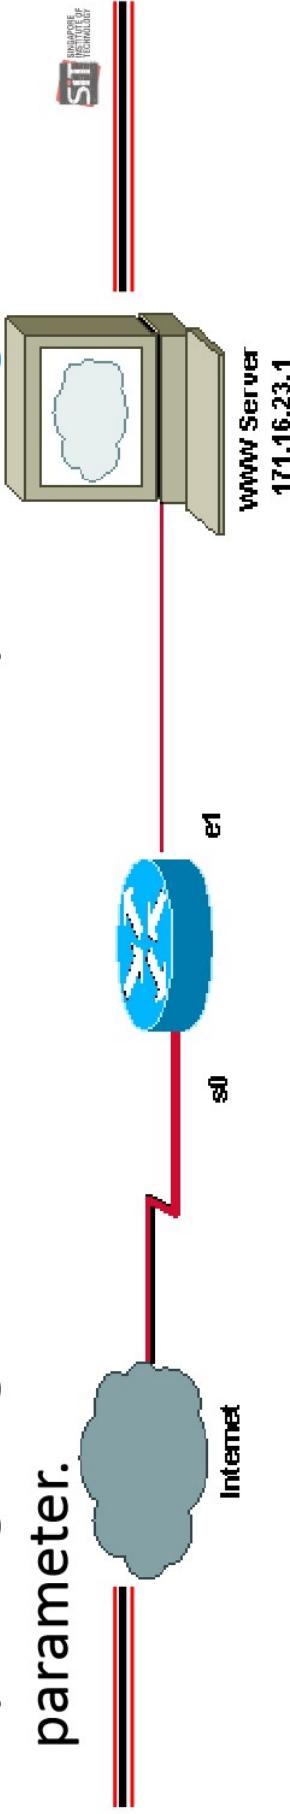
```
access-list 100 permit tcp any host 171.16.23.1 eq 80  
access-list 100 deny ip any any
```

A screenshot of Wireshark version 1.8.10 displaying a capture file named 'fragattack-def1.pcapng'. The interface is set to 'Expression...'. The packet list shows four fragments of a TCP connection. The first three fragments have off=0, ID=0219, and the fourth has off=24, ID=0219. The details and bytes panes show the fragmented nature of the TCP segments.

No.	Time	Source	Destination	Protocol	Length	Info
2	4.185729	192.168.1.10	171.16.23.1	IPv4	60	Fragmented IP protocol (proto=TCP, off=0, ID=0219)
3	4.185765	192.168.1.10	171.16.23.1	IPv4	60	Fragmented IP protocol (proto=TCP, off=16, ID=0219)
4	4.185766	192.168.1.10	171.16.23.1	IPv4	60	Fragmented IP protocol (proto=TCP, off=24, ID=0219)

Notice fragment FO=1, denoted as off=8 in Wireshark, was dropped.

Alternatively, filtering of **non-initial fragments** can also be done by configuring **ACL statements with the optional ‘fragments’ parameter.**



Applying following ACL for **ingress filtering** at interface s0:

```
access-list 101 deny tcp any host 171.16.23.1 fragments
access-list 101 permit tcp any host 171.16.23.1 eq 80
access-list 101 deny ip any any
```

Screenshot of Wireshark showing captured traffic. A blue arrow points to the 'Expression...' dropdown menu at the top right of the packet list.

Packet list table:

No.	Time	Source	Destination	Protocol	Length	Info
2	0.377960	192.168.1.10	171.16.23.1	IPv4	60	Fragmented IP protocol (proto=TCP 6, off=0, ID=01d7)
3	2.285187	192.168.1.10	171.16.23.1	IPv4	60	Fragmented IP protocol (proto=TCP 6, off=0, ID=01d9)
4	3.372892	192.168.1.10	171.16.23.1	IPv4	60	Fragmented IP protocol (proto=TCP 6, off=0, ID=01db)
5	5.291617	192.168.1.10	171.16.23.1	IPv4	60	Fragmented IP protocol (proto=TCP 6, off=0, ID=01dd)

Notice all fragments are now dropped except initial fragment.